

DATABASE MANAGEMENT SYSTEMS SOLUTIONS MANUAL THIRD EDITION

Raghu Ramakrishnan
*University of Wisconsin
Madison, WI, USA*

Johannes Gehrke
*Cornell University
Ithaca, NY, USA*

Jeff Derstadt, Scott Selikoff, and Lin Zhu
*Cornell University
Ithaca, NY, USA*

CONTENTS

PREFACE	iii	1 INTRODUCTION TO DATABASE SYSTEMS	1	2
INTRODUCTION TO DATABASE DESIGN	6	3 THE RELATIONAL		
MODEL	16	4 RELATIONAL ALGEBRA AND CALCULUS	28	5 SQL:
QUERIES, CONSTRAINTS, TRIGGERS	45	6 DATABASE		
APPLICATION DEVELOPMENT	63	7 INTERNET APPLICATIONS	66	8

OVERVIEW OF STORAGE AND INDEXING 73 9 STORING DATA:
DISKS AND FILES 81 10 TREE-STRUCTURED INDEXING 88 11
HASH-BASED INDEXING 100 12 OVERVIEW OF QUERY
EVALUATION 119 13 EXTERNAL SORTING 126 14 EVALUATION OF
RELATIONAL OPERATORS 131

i

iiDatabase Management Systems Solutions Manual Third Edition

15 A TYPICAL QUERY OPTIMIZER 144 16 OVERVIEW OF
TRANSACTION MANAGEMENT 159 17 CONCURRENCY CONTROL
167 18 CRASH RECOVERY 179 19 SCHEMA REFINEMENT AND
NORMAL FORMS 189 20 PHYSICAL DATABASE DESIGN AND
TUNING 204 21 SECURITY 215

PREFACE

It is not every question that deserves an answer.

Publius Syrus, 42 B.C.

I hope that most of the questions in this book deserve an answer. The set of questions is unusually extensive, and is designed to reinforce and deepen students' understanding of the concepts covered in each chapter. There is a strong emphasis on quantitative and problem-solving type exercises.

While I wrote some of the solutions myself, most were written originally by students in the database classes at Wisconsin. I'd like to thank the many students who helped in developing and checking the solutions to the exercises; this manual would not be available without their contributions. In alphabetical order: X. Bao, S. Biao, M. Chakrabarti, C. Chan, W. Chen, N. Cheung, D. Colwell, J. Derstadt, C. Fritz, V. Ganti, J. Gehrke, G. Glass, V. Gopalakrishnan, M. Higgins, T. Jasmin, M. Krishnaprasad, Y. Lin, C. Liu, M. Lusignan, H. Modi, S. Narayanan, D. Randolph, A. Ranganathan, J. Reminga, A. Therber, M. Thomas, Q. Wang, R. Wang, Z. Wang and J. Yuan. In addition, James Harrington and Martin Reames at Wisconsin and Nina Tang at Berkeley provided especially detailed feedback.

Several students contributed to each chapter's solutions, and answers were subse

quently checked by me and by other students. This manual has been in use for several semesters. I hope that it is now mostly accurate, but I'm sure it still contains errors and omissions. If you are a student and you do not understand a particular solution, contact your instructor; it may be that you are missing something, but it may also be that the solution is incorrect! If you discover a bug, please send me mail (raghu@cs.wisc.edu) and I will update the manual promptly.

The latest version of this solutions manual is distributed freely through the Web; go to the home page mentioned below to obtain a copy.

For More Information

The home page for this book is at URL:

iii

Database Management Systems Solutions Manual Third Edition

<http://www.cs.wisc.edu/~dbbook>

This page is frequently updated and contains information about the book, past and current users, and the software. This page also contains a link to all known errors in the book, the accompanying slides, and the software. *Since the solutions manual is distributed electronically, all known errors are immediately fixed and no list of errors is maintained.* Instructors are advised to visit this site periodically; they can also register at this site to be notified of important changes by email.

1

INTRODUCTION TO DATABASE SYSTEMS

Exercise 1.1 Why would you choose a database system instead of simply storing data in operating system files? When would it make sense *not* to use a database system?

Answer 1.1 A *database* is an integrated collection of data, usually so large that it has to be stored on secondary storage devices such as disks or tapes. This data can be maintained as a collection of operating system files, or stored in a *DBMS* (database management system). The advantages of using a DBMS are:

Data independence and efficient access. Database application programs are in

dependent of the details of data representation and storage. The conceptual and external schemas provide independence from physical storage decisions and logical design decisions respectively. In addition, a DBMS provides efficient storage and retrieval mechanisms, including support for very large files, index structures and query optimization.

Reduced application development time. Since the DBMS provides several important functions required by applications, such as concurrency control and crash recovery, high level query facilities, etc., only application-specific code needs to be written. Even this is facilitated by suites of application development tools available from vendors for many database management systems.

Data integrity and security. The view mechanism and the authorization facilities of a DBMS provide a powerful access control mechanism. Further, updates to the data that violate the semantics of the data can be detected and rejected by the DBMS if users specify the appropriate *integrity constraints*.

Data administration. By providing a common umbrella for a large collection of data that is shared by several users, a DBMS facilitates maintenance and data administration tasks. A good DBA can effectively shield end-users from the chores of fine-tuning the data representation, periodic back-ups etc.

Concurrent access and crash recovery. A DBMS supports the notion of a *transaction*, which is conceptually a single user's sequential program. Users can write transactions as if their programs were running in isolation against the database. The DBMS executes the actions of transactions in an interleaved fashion to obtain good performance, but schedules them in such a way as to ensure that conflicting operations are not permitted to proceed concurrently. Further, the DBMS maintains a continuous log of the changes to the data, and if there is a system crash, it can restore the database to a *transaction-consistent* state. That is, the actions of incomplete transactions are undone, so that the database state reflects only the actions of completed transactions. Thus, if each complete transaction, executing alone, maintains the consistency criteria, then the database state after recovery from a crash is consistent.

If these advantages are not important for the application at hand, using a collection of files may be a better solution because of the increased cost and overhead of purchasing and maintaining a DBMS.

Exercise 1.2 What is logical data independence and why is it important?

Answer 1.2 Answer omitted.

Exercise 1.3 Explain the difference between logical and physical data independence.

Answer 1.3 Logical data independence means that users are shielded from changes in the logical structure of the data, while physical data independence insulates users from changes in the physical storage of the data. We saw an example of logical data independence in the answer to Exercise 1.2. Consider the Students relation from that example (and now assume that it is not replaced by the two smaller relations). We could choose to store Students tuples in a heap file, with a clustered index on the sname field. Alternatively, we could choose to store it with an index on the gpa

field, or to create indexes on both fields, or to store it as a file sorted by gpa. These storage alternatives are not visible to users, except in terms of improved performance, since they simply see a relation as a set of tuples. This is what is meant by physical data independence.

Exercise 1.4 Explain the difference between external, internal, and conceptual schemas. How are these different schema layers related to the concepts of logical and physical data independence?

Answer 1.4 Answer omitted.

Exercise 1.5 What are the responsibilities of a DBA? If we assume that the DBA is never interested in running his or her own queries, does the DBA still need to understand query optimization? Why?

Introduction to Database Systems 3 Answer 1.5 The DBA is responsible for:

Designing the logical and physical schemas, as well as widely-used portions of the external schema.

Security and authorization.

Data availability and recovery from failures.

Database tuning: The DBA is responsible for evolving the database, in particular the conceptual and physical schemas, to ensure adequate performance as user requirements change.

A DBA needs to understand query optimization even if s/he is not interested in running his or her own queries because some of these responsibilities (database design and tuning) are related to query optimization. Unless the DBA understands the performance needs of widely used queries, and how the DBMS will optimize and execute these queries, good design and tuning decisions cannot be made.

Exercise 1.6 Scrooge McNugget wants to store information (names, addresses, descriptions of embarrassing moments, etc.) about the many ducks on his payroll. Not surprisingly, the volume of data compels him to buy a database system. To save money, he wants to buy one with the fewest possible features, and he plans to run it as a stand-alone application on his PC clone. Of course, Scrooge does not plan to share his list with anyone. Indicate which of the following DBMS features Scrooge should pay for; in each case, also indicate why Scrooge should (or should not) pay for that feature in the system he buys.

1. A security facility.
2. Concurrency control.
3. Crash recovery.
4. A view mechanism.
5. A query language.

Answer 1.6 Answer omitted.

Exercise 1.7 Which of the following plays an important role in *representing* informa

tion about the real world in a database? Explain briefly.

1. The data definition language.

4 Chapter 1

2. The data manipulation language.

3. The buffer manager.

4. The data model.

Answer 1.7 Let us discuss the choices in turn.

The data definition language is important in representing information because it is used to describe external and logical schemas.

The data manipulation language is used to access and update data; it is not important for representing the data. (Of course, the data manipulation language must be aware of how data is represented, and reflects this in the constructs that it supports.)

The buffer manager is not very important for representation because it brings arbitrary disk pages into main memory, independent of any data representation.

The data model is fundamental to representing information. The data model determines what data representation mechanisms are supported by the DBMS. The data definition language is just the specific set of language constructs available to describe an actual application's data in terms of the *data model*.

Exercise 1.8 Describe the structure of a DBMS. If your operating system is upgraded to support some new functions on OS files (e.g., the ability to force some sequence of bytes to disk), which layer(s) of the DBMS would you have to rewrite to take advantage of these new functions?

Answer 1.8 Answer omitted.

Exercise 1.9 Answer the following questions:

1. What is a transaction?
2. Why does a DBMS interleave the actions of different transactions instead of executing transactions one after the other?
3. What must a user guarantee with respect to a transaction and database consistency? What should a DBMS guarantee with respect to concurrent execution of several transactions and database consistency?
4. Explain the strict two-phase locking protocol.
5. What is the WAL property, and why is it important?

Introduction to Database Systems 5

Answer 1.9 Let us answer each question in turn:

1. A transaction is any one execution of a user program in a DBMS. This is the basic unit of change in a DBMS.

2. A DBMS is typically shared among many users. Transactions from these users can be interleaved to improve the execution time of users' queries. By interleaving queries, users do not have to wait for other user's transactions to complete fully before their own transaction begins. Without interleaving, if user A begins a transaction that will take 10 seconds to complete, and user B wants to begin a transaction, user B would have to wait an additional 10 seconds for user A's transaction to complete before the database would begin processing user B's request.
3. A user must guarantee that his or her transaction does not corrupt data or insert nonsense in the database. For example, in a banking database, a user must guarantee that a cash withdraw transaction accurately models the amount a person removes from his or her account. A database application would be worthless if a person removed 20 dollars from an ATM but the transaction set their balance to zero! A DBMS must guarantee that transactions are executed fully and independently of other transactions. An essential property of a DBMS is that a transaction should execute atomically, or as if it is the only transaction running. Also, transactions will either complete fully, or will be aborted and the database returned to its initial state. This ensures that the database remains consistent.
4. Strict two-phase locking uses shared and exclusive locks to protect data. A transaction must hold all the required locks before executing, and does not release any lock until the transaction has completely finished.
5. The WAL property affects the logging strategy in a DBMS. The WAL, Write Ahead Log, property states that each write action must be recorded in the log (on disk) before the corresponding change is reflected in the database itself. This protects the database from system crashes that happen during a transaction's execution. By recording the change in a log before the change is truly made, the database knows to undo the changes to recover from a system crash. Otherwise, if the system crashes just after making the change in the database but before the database logs the change, then the database would not be able to detect this change during crash recovery.

2

INTRODUCTION TO DATABASE DESIGN

Exercise 2.1 Explain the following terms briefly: *attribute*, *domain*, *entity*, *relationship*, *entity set*, *relationship set*, *one-to-many relationship*, *many-to-many relationship*, *participation constraint*, *overlap constraint*, *covering constraint*, *weak entity set*, *aggregation*, and *role indicator*.

Answer 2.1 Term explanations:

Attribute - a property or description of an entity. A toy department employee entity could have attributes describing the employee's name, salary, and years of service.

Domain - a set of possible values for an attribute.

Entity - an object in the real world that is distinguishable from other objects such as the green dragon toy.

Relationship - an association among two or more entities.

Entity set - a collection of similar entities such as all of the toys in the toy department.

Relationship set - a collection of similar relationships

One-to-many relationship - a key constraint that indicates that one entity can be associated with many of another entity. An example of a one-to-many relationship is when an employee can work for only one department, and a department can have many employees.

Many-to-many relationship - a key constraint that indicates that many of one entity can be associated with many of another entity. An example of a many-to-many relationship is employees and their hobbies: a person can have many different hobbies, and many people can have the same hobby.

Participation constraint - a participation constraint determines whether relationships must involve certain entities. An example is if every department entity has a manager entity. Participation constraints can either be total or partial. A total participation constraint says that every department has a manager. A partial participation constraint says that every employee does not have to be a manager.

Overlap constraint - within an ISA hierarchy, an overlap constraint determines whether or not two subclasses can contain the same entity.

Covering constraint - within an ISA hierarchy, a covering constraint determines where the entities in the subclasses collectively include all entities in the superclass. For example, with an Employees entity set with subclasses HourlyEmployee and SalaryEmployee, does every Employee entity necessarily have to be within either HourlyEmployee or SalaryEmployee?

Weak entity set - an entity that cannot be identified uniquely without considering some primary key attributes of another identifying owner entity. An example is including Dependent information for employees for insurance purposes.

Aggregation - a feature of the entity relationship model that allows a relationship set to participate in another relationship set. This is indicated on an ER diagram by drawing a dashed box around the aggregation.

Role indicator - If an entity set plays more than one role, role indicators describe the different purpose in the relationship. An example is a single Employee entity set with a relation Reports-To that relates supervisors and subordinates.

Exercise 2.2 A university database contains information about professors (identified by social security number, or SSN) and courses (identified by courseid). Professors teach courses; each of the following situations concerns the Teaches relationship set. For each situation, draw an ER diagram that describes it (assuming no further constraints hold).

1. Professors can teach the same course in several semesters, and each offering must be recorded.
2. Professors can teach the same course in several semesters, and only the most recent such offering needs to be recorded. (Assume this condition applies in all subsequent questions.)
3. Every professor must teach some course.
4. Every professor teaches exactly one course (no more, no less).
5. Every professor teaches exactly one course (no more, no less), and every course must be taught by some professor.

8 Chapter 2

6. Now suppose that certain courses can be taught by a team of professors jointly, but it is possible that no one professor in a team can teach the course. Model this situation, introducing additional entity sets and relationship sets if necessary.

Answer 2.2 Answer omitted.

Exercise 2.3 Consider the following information about a university database:

Professors have an SSN, a name, an age, a rank, and a research specialty.

Projects have a project number, a sponsor name (e.g., NSF), a starting date, an ending date, and a budget.

Graduate students have an SSN, a name, an age, and a degree program (e.g., M.S. or Ph.D.).

Each project is managed by one professor (known as the project's principal investigator).

Each project is worked on by one or more professors (known as the project's co-investigators).

Professors can manage and/or work on multiple projects.

Each project is worked on by one or more graduate students (known as the project's research assistants).

When graduate students work on a project, a professor must supervise their work on the project. Graduate students can work on multiple projects, in which case they will have a (potentially different) supervisor for each one.

Departments have a department number, a department name, and a main office. Departments have a professor (known as the chairman) who runs the

department.

Professors work in one or more departments, and for each department that they work in, a time percentage is associated with their job.

Graduate students have one major department in which they are working on their degree.

Each graduate student has another, more senior graduate student (known as a student advisor) who advises him or her on what courses to take.

Design and draw an ER diagram that captures the information about the university. Use only the basic ER model here; that is, entities, relationships, and attributes. Be sure to indicate any key and participation constraints.

41
4
4
4
4
4

41
4
4
4
4
4
4
4
4
4
4

4
4
4
4
4
4
4
4
4
4

4
4
4
4
4
4
4
4
4
4
4

41
4
4
4
4
4
4
4
4
4
4

4
4

Figure 2.1 ER Diagram for Exercise 2.3

10 Chapter 2 Answer 2.3 The ER diagram is shown in Figure 2.1.

Exercise 2.4 A company database needs to store information about employees (identified by *ssn*, with *salary* and *phone* as attributes), departments (identified by *dno*, with *dname* and *budget* as attributes), and children of employees (with *name* and *age* as attributes). Employees *work* in departments; each department is *managed by* an employee; a child must be identified uniquely by *name* when the parent (who is an employee; assume that only one parent works for the company) is known. We are not interested in information about a child once the parent leaves the company.

Draw an ER diagram that captures this information.

Answer 2.4 Answer omitted.

Exercise 2.5 Notown Records has decided to store information about musicians who perform on its albums (as well as other company data) in a database. The company has wisely chosen to hire you as a database designer (at your usual consulting fee of \$2500/day).

Each musician that records at Notown has an SSN, a name, an address, and a phone number. Poorly paid musicians often share the same address, and no address has more than one phone.

Introduction to Database Design 11

Figure 2.2 ER Diagram for Exercise 2.5

Design a conceptual schema for Notown and draw an ER diagram for your schema. The preceding information describes the situation that the Notown database must model. Be sure to indicate all key and cardinality constraints and any assumptions you make. Identify any constraints you are unable to capture in the ER diagram and briefly explain why you could not express them.

Answer 2.5 The ER diagram is shown in Figure 2.2.

12 Chapter 2

Exercise 2.6 Computer Sciences Department frequent fliers have been complaining to Dane County Airport officials about the poor organization at the airport. As a result, the officials decided that all information related to the airport should be organized using a DBMS, and you have been hired to design the database. Your first task is to organize the information about all the airplanes stationed and maintained at the airport. The relevant information is as follows:

Every airplane has a registration number, and each airplane is of a specific model.

The airport accommodates a number of airplane models, and each model is identified by a model number (e.g., DC-10) and has a capacity and a weight.

A number of technicians work at the airport. You need to store the name, SSN, address, phone number, and salary of each technician.

Each technician is an expert on one or more plane model(s), and his or her expertise may overlap with that of other technicians. This information about technicians must also be recorded.

Traffic controllers must have an annual medical examination. For each traffic controller, you must store the date of the most recent exam.

All airport employees (including technicians) belong to a union. You must store the union membership number of each employee. You can assume that each employee is uniquely identified by a social security number.

The airport has a number of tests that are used periodically to ensure that airplanes are still airworthy. Each test has a Federal Aviation Administration (FAA) test number, a name, and a maximum possible score.

The FAA requires the airport to keep track of each time a given airplane is tested by a given technician using a given test. For each testing event, the information needed is the date, the number of hours the technician spent doing the test, and the score the airplane received on the test.

1. Draw an ER diagram for the airport database. Be sure to indicate the various attributes of each entity and relationship set; also specify the key and participation constraints for each relationship set. Specify any necessary

overlap and covering constraints as well (in English).

2. The FAA passes a regulation that tests on a plane must be conducted by a technician who is an expert on that model. How would you express this constraint in the ER diagram? If you cannot express it, explain briefly.

Answer 2.6 Answer omitted.

Introduction to Database Design 13

Exercise 2.7 The Prescriptions-R-X chain of pharmacies has offered to give you a free lifetime supply of medicine if you design its database. Given the rising cost of health care, you agree. Here's the information that you gather:

Patients are identified by an SSN, and their names, addresses, and ages must be recorded.

Doctors are identified by an SSN. For each doctor, the name, specialty, and years of experience must be recorded.

Each pharmaceutical company is identified by name and has a phone number.

For each drug, the trade name and formula must be recorded. Each drug is sold by a given pharmaceutical company, and the trade name identifies a drug uniquely from among the products of that company. If a pharmaceutical company is deleted, you need not keep track of its products any longer.

Each pharmacy has a name, address, and phone number.

Every patient has a primary physician. Every doctor has at least one patient.

Each pharmacy sells several drugs and has a price for each. A drug could be sold at several pharmacies, and the price could vary from one pharmacy to another.

Doctors prescribe drugs for patients. A doctor could prescribe one or more drugs for several patients, and a patient could obtain prescriptions from several doctors. Each prescription has a date and a quantity associated with it. You can assume that, if a doctor prescribes the same drug for the same patient more than once, only the last such prescription needs to be stored.

Pharmaceutical companies have long-term contracts with pharmacies. A pharmaceutical company can contract with several pharmacies, and a pharmacy can contract with several pharmaceutical companies. For each contract, you have to store a start date, an end date, and the text of the contract.

Pharmacies appoint a supervisor for each contract. There must always be a supervisor for each contract, but the contract supervisor can change over the lifetime of the contract.

1. Draw an ER diagram that captures the preceding information. Identify any constraints not captured by the ER diagram.
2. How would your design change if each drug must be sold at a fixed price by all pharmacies?
3. How would your design change if the design requirements change as follows: If a doctor prescribes the same drug for the same patient more than once, several such prescriptions may have to be stored.

Answer 2.8 Answer omitted.

Exercise 2.9 Answer the following questions.

Explain the following terms briefly: *UML*, *use case diagrams*, *statechart diagrams*, *class diagrams*, *database diagrams*, *component diagrams*, and *deployment diagrams*.

Explain the relationship between ER diagrams and UML.

Answer 2.9 Not yet done.

3

THE RELATIONAL MODEL

Exercise 3.1 Define the following terms: *relation schema*, *relational database schema*, *domain*, *attribute*, *attribute domain*, *relation instance*, *relation cardinality*, and *relation degree*.

Answer 3.1 A *relation schema* can be thought of as the basic information describing a table or *relation*. This includes a set of column names, the data types associated with each column, and the name associated with the entire table. For example, a relation schema for the relation called Students could be expressed using the following representation:

Students(*sid*: string, *name*: string, *login*: string,
age: integer, *gpa*: real)

There are five fields or columns, with names and types as shown above.

A *relational database schema* is a collection of relation schemas, describing one or more relations.

Domain is synonymous with *data type*. *Attributes* can be thought of as columns in a table. Therefore, an *attribute domain* refers to the data type associated with a column.

A *relation instance* is a set of tuples (also known as *rows* or *records*) that each conform to the schema of the relation.

The *relation cardinality* is the number of tuples in the relation.

The *relation degree* is the number of fields (or columns) in the relation.

Exercise 3.2 How many distinct tuples are in a relation instance with cardinality 22?

Answer 3.2 Answer omitted.

16

The Relational Model 17

Exercise 3.3 Does the relational model, as seen by an SQL query writer, provide physical and logical data independence? Explain.

Answer 3.3 The user of SQL has no idea how the data is physically represented in the machine. He or she relies entirely on the relation abstraction for querying. Physical data independence is therefore assured. Since a user can define views, logical data independence can also be achieved by using view definitions to hide changes in the conceptual schema.

Exercise 3.4 What is the difference between a candidate key and the primary key for a given relation? What is a superkey?

Answer 3.4 Answer omitted.

FIELDS (ATTRIBUTES, COLUMNS)									
<i>sid age gpa name login 50000 3.3</i>									
Field names									
TUPLES (RECORDS, ROWS)	53666 53688 53650 53831 53832 Dave	Jones	Guldu	smith@math	19	18	18	19	
		Smith	dave@cs	madayan@m	11	12			
		Smith	jones@cs	usic	3.4	3.2	3.8		
		Madayan	smith@ee	guldu@music	1.8	2.0			

Figure 3.1 An Instance S1 of the Students Relation

Exercise 3.5 Consider the instance of the Students relation shown in Figure 3.1.

1. Give an example of an attribute (or set of attributes) that you can deduce is *not* a candidate key, based on this instance being legal.
2. Is there any example of an attribute (or set of attributes) that you can deduce is a candidate key, based on this instance being legal?

Answer 3.5 Examples of non-candidate keys include the following: {name}, {age}. (Note that {gpa} can *not* be declared as a non-candidate key from this evidence alone even though common sense tells us that clearly more than one student could have the same grade point average.)

You cannot determine a key of a relation given only one instance of the relation. The fact that the instance is “legal” is immaterial. A candidate key, as defined here, is a

key, not something that only *might* be a key. The instance shown is just one possible

“snapshot” of the relation. At other times, the same relation may have an instance (or snapshot) that contains a totally different set of tuples, and we cannot make predictions about those instances based only upon the instance that we are given.

Exercise 3.6 What is a foreign key constraint? Why are such constraints important? What is referential integrity?

Answer 3.6 Answer omitted.

Exercise 3.7 Consider the relations Students, Faculty, Courses, Rooms, Enrolled, Teaches, and Meets In defined in Section 1.5.2.

1. List all the foreign key constraints among these relations.
2. Give an example of a (plausible) constraint involving one or more of these relations that is not a primary key or foreign key constraint.

Answer 3.7 There is no reason for a foreign key constraint (FKC) on the Students, Faculty, Courses, or Rooms relations. These are the most basic relations and must be free-standing. Special care must be given to entering data into these base relations.

In the Enrolled relation, *sid* and *cid* should both have FKCs placed on them. (Real students must be enrolled in real courses.) Also, since real teachers must teach real courses, both the *f id* and the *cid* fields in the Teaches relation should have FKCs. Finally, Meets In should place FKCs on both the *cid* and *mo* fields.

It would probably be wise to enforce a few other constraints on this DBMS: the length of *sid*, *cid*, and *f id* could be standardized; checksums could be added to these identification numbers; limits could be placed on the size of the numbers entered into the credits, capacity, and salary fields; an enumerated type should be assigned to the grade field (preventing a student from receiving a grade of G, among other things); etc.

Exercise 3.8 Answer each of the following questions briefly. The questions are based on the following relational schema:

Emp(*eid*: integer, *ename*: string, *age*: integer, *salary*: real)
Works(*eid*: integer, *did*: integer, *pctime*: integer)
Dept(*did*: integer, *dname*: string, *budget*: real, *managerid*: integer)

1. Give an example of a foreign key constraint that involves the Dept relation. What are the options for enforcing this constraint when a user attempts to delete a Dept tuple?

The Relational Model 19

2. Write the SQL statements required to create the preceding relations, including appropriate versions of all primary and foreign key integrity constraints.
3. Define the Dept relation in SQL so that every department is guaranteed to have a manager.
4. Write an SQL statement to add John Doe as an employee with *eid* = 101, *age* = 32 and *salary* = 15, 000.

5. Write an SQL statement to give every employee a 10 percent raise.
6. Write an SQL statement to delete the Toy department. Given the referential integrity constraints you chose for this schema, explain what happens when this statement is executed.

Answer 3.8 Answer omitted.

	<i>sid</i>	<i>name</i>	<i>login</i>	<i>age</i>	<i>gpa</i>
	53831	Madayan	madayan@music	11	1.8
	53832	Guldu	guldu@music	12	2.0

Figure 3.2 Students with *age* < 18 on Instance S

Exercise 3.9 Consider the SQL query whose answer is shown in Figure 3.2.

1. Modify this query so that only the *login* column is included in the answer.
2. If the clause WHERE *S.gpa* >= 2 is added to the original query, what is the set of tuples in the answer?

Answer 3.9 The answers are as follows:

1. Only *login* is included in the answer:

```
SELECT S.login
FROM Students S
WHERE S.age < 18
```

2. The answer tuple for Madayan is omitted then.

Exercise 3.10 Explain why the addition of NOT NULL constraints to the SQL definition of the Manages relation (in Section 3.5.3) does not enforce the constraint that each department must have a manager. What, if anything, is achieved by requiring that the *ssn* field of Manages be non-*null*?

20 Chapter 3 Answer 3.10 Answer omitted.

Exercise 3.11 Suppose that we have a ternary relationship R between entity sets A, B, and C such that A has a key constraint and total participation and B has a key constraint; these are the only constraints. A has attributes *a1* and *a2*, with *a1* being the key; B and C are similar. R has no descriptive attributes. Write SQL statements that create tables corresponding to this information so as to capture as many of the constraints as possible. If you cannot capture some constraint, explain why.

Answer 3.11 The following SQL statements create the corresponding relations.

```
CREATE TABLE A ( a1 CHAR(10),
                  a2 CHAR(10),
                  b1 CHAR(10),
                  c1 CHAR(10),
                  PRIMARY KEY (a1),
                  UNIQUE (b1),
                  FOREIGN KEY (b1) REFERENCES B,
```


FOREIGN KEY (prof ssn) REFERENCES Professors,
FOREIGN KEY (dno) REFERENCES Depts)

Observe that we would need check constraints or assertions in SQL to enforce the rule that Professors work in at least one department.

5. CREATE TABLE Project (pid INTEGER,
 sponsor CHAR(32),
 start date DATE,
 end date DATE,
 budget FLOAT,
 PRIMARY KEY (pid))

6. CREATE TABLE Graduates (grad ssn CHAR(10),
 age INTEGER,
 name CHAR(64),
 deg prog CHAR(32),

22 Chapter 3

 major INTEGER,
 PRIMARY KEY (grad ssn),
 FOREIGN KEY (major) REFERENCES Depts)

Note that the Major table is not necessary since each Graduate has only one major and so this can be an attribute in the Graduates table.

7. CREATE TABLE Advisor (senior ssn CHAR(10),
 grad ssn CHAR(10),
 PRIMARY KEY (senior ssn, grad ssn),
 FOREIGN KEY (senior ssn)
 REFERENCES Graduates (grad ssn),
 FOREIGN KEY (grad ssn) REFERENCES Graduates)

8. CREATE TABLE Manages (pid INTEGER,
 prof ssn CHAR(10),
 PRIMARY KEY (pid, prof ssn),
 FOREIGN KEY (prof ssn) REFERENCES Professors,
 FOREIGN KEY (pid) REFERENCES Projects)

9. CREATE TABLE Work In (pid INTEGER,
 prof ssn CHAR(10),
 PRIMARY KEY (pid, prof ssn),
 FOREIGN KEY (prof ssn) REFERENCES Professors,
 FOREIGN KEY (pid) REFERENCES Projects)

Observe that we cannot enforce the participation constraint for Projects in the Work In table without check constraints or assertions in SQL.

10. CREATE TABLE Supervises (prof ssn CHAR(10),
 grad ssn CHAR(10),
 pid INTEGER,
 PRIMARY KEY (prof ssn, grad ssn, pid),
 FOREIGN KEY (prof ssn) REFERENCES Professors,
 FOREIGN KEY (grad ssn) REFERENCES Graduates,

FOREIGN KEY (pid) REFERENCES Projects)

Note that we do not need an explicit table for the Work Proj relation since every time a Graduate works on a Project, he or she must have a Supervisor.

Exercise 3.14 Consider the scenario from Exercise 2.4, where you designed an ER diagram for a company database. Write SQL statements to create the corresponding
The Relational Model 23

relations and capture as many of the constraints as possible. If you cannot capture some constraints, explain why.

Answer 3.14 Answer omitted.

Exercise 3.15 Consider the Notown database from Exercise 2.5. You have decided to recommend that Notown use a relational database system to store company data. Show the SQL statements for creating relations corresponding to the entity sets and relationship sets in your design. Identify any constraints in the ER diagram that you are unable to capture in the SQL statements and briefly explain why you could not express them.

Answer 3.15 The following SQL statements create the corresponding relations.

1. CREATE TABLE Musicians (ssn CHAR(10),
 name CHAR(30),
 PRIMARY KEY (ssn))
2. CREATE TABLE Instruments (instrId CHAR(10),
 dname CHAR(30),
 key CHAR(5),
 PRIMARY KEY (instrId))
3. CREATE TABLE Plays (ssn CHAR(10),
 instrId INTEGER,
 PRIMARY KEY (ssn, instrId),
 FOREIGN KEY (ssn) REFERENCES Musicians,
 FOREIGN KEY (instrId) REFERENCES Instruments)
4. CREATE TABLE Songs Appears (songId INTEGER,
 author CHAR(30),
 title CHAR(30),
 albumIdentifier INTEGER NOT NULL,
 PRIMARY KEY (songId),
 FOREIGN KEY (albumIdentifier)
 References Album Producer,
5. CREATE TABLE Telephone Home (phone CHAR(11),
 address CHAR(30),
 PRIMARY KEY (phone),
 FOREIGN KEY (address) REFERENCES Place,

```
1. CREATE TABLE Pri Phy Patient ( ssn CHAR(11),
                                   name CHAR(20),
                                   age INTEGER,
                                   address CHAR(20),
```

```

phy ssn CHAR(11),
PRIMARY KEY (ssn),
FOREIGN KEY (phy ssn) REFERENCES Doctor )

```

```

2. CREATE TABLE Prescription ( ssn CHAR(11),
    phy ssn CHAR(11),
    date CHAR(11),
    quantity INTEGER,
    trade name CHAR(20),
    pharm id CHAR(11),
    PRIMARY KEY (ssn, phy ssn),
    FOREIGN KEY (ssn) REFERENCES Patient,
    FOREIGN KEY (phy ssn) REFERENCES Doctor,
    FOREIGN KEY (trade name, pharm id)
    References Make Drug)

```

```

3. CREATE TABLE Make Drug (trade name CHAR(20),
    pharm id CHAR(11),
    PRIMARY KEY (trade name, pharm id),
    FOREIGN KEY (trade name) REFERENCES Drug,
    FOREIGN KEY (pharm id) REFERENCES Pharm co)

```

```

4. CREATE TABLE Sell ( price INTEGER,
    name CHAR(10),
    trade name CHAR(10),
    PRIMARY KEY (name, trade name),
    FOREIGN KEY (name) REFERENCES Pharmacy,
    FOREIGN KEY (trade name) REFERENCES Drug)

```

```

5. CREATE TABLE Contract ( name CHAR(20),
    pharm id CHAR(11),
    start date CHAR(11),
    end date CHAR(11),

```

26 Chapter 3

```

text CHAR(10000),
supervisor CHAR(20),
PRIMARY KEY (name, pharm id),
FOREIGN KEY (name) REFERENCES Pharmacy,
FOREIGN KEY (pharm id) REFERENCES Pharm co)

```

Exercise 3.18 Write SQL statements to create the corresponding relations to the ER diagram you designed for Exercise 2.8. If your translation cannot capture any constraints in the ER diagram, explain why.

Answer 3.18 Answer omitted.

Exercise 3.19 Briefly answer the following questions based on this schema:

Emp(*eid*: integer, *ename*: string, *age*: integer, *salary*: real)

Works(*eid*: integer, *did*: integer, *pct time*: integer)
Dept(*did*: integer, *budget*: real, *managerid*: integer)

1. Suppose you have a view SeniorEmp defined as follows:

```
CREATE VIEW SeniorEmp (sname, sage, salary)
AS SELECT E.ename, E.age, E.salary
FROM Emp E
WHERE E.age > 50
```

Explain what the system will do to process the following query:

```
SELECT S.sname
FROM SeniorEmp S
WHERE S.salary > 100,000
```

2. Give an example of a view on Emp that could be automatically updated by updating Emp.
3. Give an example of a view on Emp that would be impossible to update (automatically) and explain why your example presents the update problem that it does.

Answer 3.19 The answer to each question is given below.

1. The system will do the following:

The Relational Model 27

```
SELECT S.name
FROM ( SELECT E.ename AS name, E.age, E.salary
FROM Emp E
WHERE E.age > 50 ) AS S
WHERE S.salary > 100000
```

2. The following view on Emp can be updated automatically by updating Emp:

```
CREATE VIEW SeniorEmp (eid, name, age, salary)
AS SELECT E.eid, E.ename, E.age, E.salary
FROM Emp E
WHERE E.age > 50
```

3. The following view cannot be updated automatically because it is not clear which employee records will be affected by a given update:

```
CREATE VIEW AvgSalaryByAge (age, avgSalary)
AS SELECT E.eid, AVG (E.salary)
FROM Emp E
GROUP BY E.age
```

Exercise 3.20 Consider the following schema:

Suppliers(*sid*: integer, *sname*: string, *address*: string)

Parts(*pid*: integer, *pname*: string, *color*: string)
Catalog(*sid*: integer, *pid*: integer, *cost*: real)

The Catalog relation lists the prices charged for parts by Suppliers. Answer the following questions:

Give an example of an updatable view involving one relation.

Give an example of an updatable view involving two relations.

Give an example of an insertable-into view that is updatable.

Give an example of an insertable-into view that is not updatable.

Answer 3.20 Answer omitted.

4

RELATIONAL ALGEBRA AND CALCULUS

Exercise 4.1 Explain the statement that relational algebra operators can be *composed*. Why is the ability to compose operators important?

Answer 4.1 Every operator in relational algebra accepts one or more relation instances as arguments and the result is always an relation instance. So the argument of one operator could be the result of another operator. This is important because, this makes it easy to write complex queries by simply composing the relational algebra operators.

Exercise 4.2 Given two relations $R1$ and $R2$, where $R1$ contains $N1$ tuples, $R2$ contains $N2$ tuples, and $N2 > N1 > 0$, give the minimum and maximum possible sizes (in tuples) for the resulting relation produced by each of the following relational algebra expressions. In each case, state any assumptions about the schemas for $R1$ and $R2$ needed to make the expression meaningful:

(1) $R1 \cup R2$, (2) $R1 \cap R2$, (3) $R1 - R2$, (4) $R1 \times R2$, (5) $\sigma_{a=5}(R1)$, (6) $\pi_a(R1)$,
and (7) $R1/R2$

Answer 4.2 Answer omitted.

Exercise 4.3 Consider the following schema:

Suppliers(sid: integer, sname: string, address: string)
 Parts(pid: integer, pname: string, color: string)
 Catalog(sid: integer, pid: integer, cost: real)

The key fields are underlined, and the domain of each field is listed after the field name. Therefore *sid* is the key for Suppliers, *pid* is the key for Parts, and *sid* and *pid* together form the key for Catalog. The Catalog relation lists the prices charged for parts by Suppliers. Write the following queries in relational algebra, tuple relational calculus, and domain relational calculus:

28

Relational Algebra and Calculus 29

1. Find the *names* of suppliers who supply some red part.
2. Find the *sids* of suppliers who supply some red or green part.
3. Find the *sids* of suppliers who supply some red part or are at 221 Packer Street.
4. Find the *sids* of suppliers who supply some red part and some green part.
5. Find the *sids* of suppliers who supply every part.
6. Find the *sids* of suppliers who supply every red part.
7. Find the *sids* of suppliers who supply every red or green part.
8. Find the *sids* of suppliers who supply every red part or supply every green part.
9. Find pairs of *sids* such that the supplier with the first *sid* charges more for some part than the supplier with the second *sid*.
10. Find the *pids* of parts supplied by at least two different suppliers.
11. Find the *pids* of the most expensive parts supplied by suppliers named Yosemite Sham.
12. Find the *pids* of parts supplied by every supplier at less than \$200. (If any supplier either does not supply the part or charges more than \$200 for it, the part is not selected.)

Answer 4.3 In the answers below RA refers to Relational Algebra, TRC refers to Tuple Relational Calculus and DRC refers to Domain Relational Calculus.

1. RA

$$\pi_{sname}(\pi_{sid}((\pi_{pid}\sigma_{color=red} P\text{ parts}) \text{ Catalog}) \text{ Suppliers}))$$

TRC

$$\{T \mid \exists T1 \in Suppliers(\exists X \in P\text{ parts}(X.color = red \wedge \exists Y \in Catalog(Y.pid = X.pid \wedge Y.sid = T1.sid)) \wedge T.sname = T1.sname)\}$$

DRC

$$\{Y \mid X, Y, Z \in Suppliers \wedge \exists P, Q, R(P, Q, R \in P\text{ parts} \wedge R = red \wedge \exists I, J, K(I, J, K \in Catalog \wedge J = P \wedge I = X))\}$$

SQL

30 Chapter 4

```
SELECT S.sname
FROM Suppliers S, Parts P, Catalog C
WHERE P.color='red' AND C.pid=P.pid AND C.sid=S.sid
```

2. RA

$$\pi_{sid}(\pi_{pid}(\sigma_{color=red \vee color=green} P \text{ parts}) \text{ catalog})$$

TRC

$$\{T \mid \exists T1 \in \text{Catalog} (\exists X \in P \text{ arts} ((X.color = 'red' \vee X.color = 'green') \wedge X.pid = T1.pid) \wedge T.sid = T1.sid)\}$$

DRC

$$\{X \mid X, Y, Z \in \text{Catalog} \wedge \exists A, B, C (A, B, C \in P \text{ arts} \wedge (C = red \vee C = green) \wedge A = Y)\}$$

SQL

```
SELECT C.sid
FROM Catalog C, Parts P
WHERE (P.color = 'red' OR P.color = 'green')
AND P.pid = C.pid
```

3. RA

$$\rho(R1, \pi_{sid}((\pi_{pid} \sigma_{color=red} P \text{ arts}) \text{ Catalog}))$$
$$\rho(R2, \pi_{sid} \sigma_{address=221PackerStreet} Suppliers)$$
$$R1 \cup R2$$

TRC

$$\{T \mid \exists T1 \in \text{Catalog} (\exists X \in P \text{ arts} (X.color = 'red' \wedge X.pid = T1.pid) \wedge T.sid = T1.sid) \\ \vee \exists T2 \in \text{Suppliers} (T2.address = 221PackerStreet \wedge T.sid = T2.sid)\}$$

DRC

$$\{X \mid X, Y, Z \in \text{Catalog} \wedge \exists A, B, C (A, B, C \in P \text{ arts} \wedge C = red \wedge A = Y) \\ \vee \exists P, Q (X, P, Q \in \text{Suppliers} \wedge Q = 221PackerStreet)\}$$

SQL

Relational Algebra and Calculus 31

```
SELECT S.sid
FROM Suppliers S
WHERE S.address = '221 Packer street'
OR S.sid IN ( SELECT C.sid
FROM Parts P, Catalog C
```


8. RA

$$\begin{aligned} & \rho(R1, ((\pi_{sid, pid} Catalog) / (\pi_{pid} \sigma_{color=red} P arts))) \\ & \rho(R2, ((\pi_{sid, pid} Catalog) / (\pi_{pid} \sigma_{color=green} P arts))) \\ & R1 \cup R2 \end{aligned}$$

TRC

$$\begin{aligned} & \{T \mid \exists T1 \in Catalog ((\forall X \in P arts \\ & (X.color = 'red \vee \exists Y \in Catalog (Y.pid = X.pid \wedge Y.sid = T \\ & 1.sid)) \vee \forall Z \in P arts (Z.color = 'green \vee \exists P \in Catalog \\ & (P.pid = Z.pid \wedge P.sid = T1.sid))) \wedge T.sid = T1.sid)\} \end{aligned}$$

DRC

$$\begin{aligned} & \{X \mid X, Y, Z \in Catalog \wedge (\forall A, B, C \in P arts \\ & (C = 'red \vee \exists P, Q, R \in Catalog (Q = A \wedge P = X)) \\ & \vee \forall U, V, W \in P arts (W = 'green \vee M, N, L \in Catalog \\ & (N = U \wedge M = X)))\} \end{aligned}$$

34 Chapter 4

SQL

```
SELECT C.sid
FROM Catalog C
WHERE (NOT EXISTS (SELECT P.pid
                    FROM Parts P
                    WHERE P.color = 'red' AND
                    (NOT EXISTS (SELECT C1.sid
                                FROM Catalog C1
                                WHERE C1.sid = C.sid AND
                                C1.pid = P.pid))))
OR ( NOT EXISTS (SELECT P1.pid
                FROM Parts P1
                WHERE P1.color = 'green' AND
                (NOT EXISTS (SELECT C2.sid
                            FROM Catalog C2
                            WHERE C2.sid = C.sid AND
                            C2.pid = P1.pid))))
```

9. RA

$$\begin{aligned} & \rho(R1, Catalog) \\ & \rho(R2, Catalog) \\ & \pi_{R1.sid, R2.sid}(\sigma_{R1.pid=R2.pid \wedge R1.sid=R2.sid \wedge R1.cost > R2.cost}(R1 \times R2)) \text{ TRC} \\ & \{T \mid \exists T1 \in Catalog (\exists T2 \in Catalog \\ & (T2.pid = T1.pid \wedge T2.sid = T1.sid \\ & \wedge T2.cost < T1.cost \wedge T.sid2 = T2.sid) \\ & \wedge T.sid1 = T1.sid)\} \end{aligned}$$

DRC

$$\{X, P \mid X, Y, Z \in \text{Catalog} \wedge \exists P, Q, R \\ (P, Q, R \in \text{Catalog} \wedge Q = Y \wedge P = X \wedge R < Z)\}$$

SQL

```
SELECT C1.sid, C2.sid
FROM Catalog C1, Catalog C2
WHERE C1.pid = C2.pid AND C1.sid = C2.sid
      AND C1.cost > C2.cost
```

Relational Algebra and Calculus 35

10. RA

$$\rho(R1, \text{Catalog}) \\ \rho(R2, \text{Catalog}) \\ \pi_{R1.pid} \sigma_{R1.pid=R2.pid \wedge R1.sid=R2.sid} (R1 \times R2)$$

TRC

$$\{T \mid \exists T1 \in \text{Catalog} (\exists T2 \in \text{Catalog} \\ (T2.pid = T1.pid \wedge T2.sid = T1.sid) \\ \wedge T.pid = T1.pid)\}$$

DRC

$$\{X \mid X, Y, Z \in \text{Catalog} \wedge \exists A, B, C \\ (A, B, C \in \text{Catalog} \wedge B = Y \wedge A = X)\}$$

SQL

```
SELECT C.pid
FROM Catalog C
WHERE EXISTS (SELECT C1.sid
              FROM Catalog C1
              WHERE C1.pid = C.pid AND C1.sid = C.sid )
```

11. RA

$$\rho(R1, \pi_{sid} \sigma_{sname=YosemiteSham} Suppliers) \\ \rho(R2, R1 \text{ Catalog}) \\ \rho(R3, R2) \\ \rho(R4(1 \rightarrow sid, 2 \rightarrow pid, 3 \rightarrow cost), \sigma_{R3.cost < R2.cost} (R3 \times R2)) \\ \pi_{pid} (R2 - \pi_{sid, pid, cost} R4)$$

TRC

$$\{T \mid \exists T1 \in \text{Catalog} (\exists X \in Suppliers \\ (X.sname = YosemiteSham \wedge X.sid = T1.sid) \wedge \neg (\exists S \in Suppliers \\ (S.sname = YosemiteSham \wedge \exists Z \in Catalog \\ (Z.sid = S.sid \wedge Z.cost > T1.cost)))) \wedge T.pid = T1.pid)\}$$

DRC

$$\{Y \mid X, Y, Z \in \text{Catalog} \wedge \exists A, B, C \\ (A, B, C \in Suppliers \wedge C = YosemiteSham \wedge A = X)\}$$

$$\wedge \neg(\exists P, Q, R(P, Q, R \in Suppliers \wedge R = YosemiteSham \\ \wedge \exists I, J, K(I, J, K \in Catalog(I = P \wedge K > Z))))\}$$

36 Chapter 4

SQL

```
SELECT C.pid
FROM Catalog C, Suppliers S
WHERE S.sname = 'Yosemite Sham' AND C.sid = S.sid
      AND C.cost ≥ ALL (Select C2.cost
                        FROM Catalog C2, Suppliers S2
                        WHERE S2.sname = 'Yosemite Sham'
                        AND C2.sid = S2.sid)
```

Exercise 4.4 Consider the Supplier-Parts-Catalog schema from the previous question. State what the following queries compute:

1. $\pi_{sname}(\pi_{sid}((\sigma_{color=red} P arts) (\sigma_{cost<100} Catalog) Suppliers))$ 2.
- $\pi_{sname}(\pi_{sid}((\sigma_{color=red} P arts) (\sigma_{cost<100} Catalog) Suppliers))$ 3.
- $(\pi_{sname}((\sigma_{color=red} P arts) (\sigma_{cost<100} Catalog) Suppliers)) \cap (\pi_{sname}((\sigma_{color=green} P arts) (\sigma_{cost<100} Catalog) Suppliers))$
4. $(\pi_{sid}((\sigma_{color=red} P arts) (\sigma_{cost<100} Catalog) Suppliers)) \cap (\pi_{sid}((\sigma_{color=green} P arts) (\sigma_{cost<100} Catalog) Suppliers))$
5. $\pi_{sname}((\pi_{sid,sname}((\sigma_{color=red} P arts) (\sigma_{cost<100} Catalog) Suppliers)) \cap (\pi_{sid,sname}((\sigma_{color=green} P arts) (\sigma_{cost<100} Catalog) Suppliers))))$

Answer 4.4 The statements can be interpreted as:

1. Find the Supplier names of the suppliers who supply a red part that costs less than 100 dollars.
2. This Relational Algebra statement does not return anything because of the sequence of projection operators. Once the sid is projected, it is the only field in the set. Therefore, projecting on sname will not return anything.
3. Find the Supplier names of the suppliers who supply a red part that costs less than 100 dollars and a green part that costs less than 100 dollars.
4. Find the Supplier ids of the suppliers who supply a red part that costs less than 100 dollars and a green part that costs less than 100 dollars.
5. Find the Supplier names of the suppliers who supply a red part that costs less than 100 dollars and a green part that costs less than 100 dollars.

Exercise 4.5 Consider the following relations containing airline flight information:

Flights(*flno*: integer, *from*: string, *to*: string,
distance: integer, *departs*: time, *arrives*: time)
Aircraft(*aid*: integer, *aname*: string, *cruisingrange*: integer)
Certified(*eid*: integer, *aid*: integer)
Employees(*eid*: integer, *ename*: string, *salary*: integer)

Note that the Employees relation describes pilots and other kinds of employees as well; every pilot is certified for some aircraft (otherwise, he or she would not qualify as a pilot), and only pilots are certified to fly.

Write the following queries in relational algebra, tuple relational calculus, and domain relational calculus. Note that some of these queries may not be expressible in relational algebra (and, therefore, also not expressible in tuple and domain relational calculus)! For such queries, informally explain why they cannot be expressed. (See the exercises at the end of Chapter 5 for additional queries over the airline schema.)

1. Find the *eids* of pilots certified for some Boeing aircraft.
2. Find the *names* of pilots certified for some Boeing aircraft.
3. Find the *aids* of all aircraft that can be used on non-stop flights from Bonn to Madras.
4. Identify the flights that can be piloted by every pilot whose salary is more than \$100,000.
5. Find the names of pilots who can operate planes with a range greater than 3,000 miles but are not certified on any Boeing aircraft.
6. Find the *eids* of employees who make the highest salary.
7. Find the *eids* of employees who make the second highest salary.
8. Find the *eids* of employees who are certified for the largest number of aircraft.
9. Find the *eids* of employees who are certified for exactly three aircraft.
10. Find the total amount paid to employees as salaries.
11. Is there a sequence of flights from Madison to Timbuktu? Each flight in the sequence is required to depart from the city that is the destination of the previous flight; the first flight must leave Madison, the last flight must reach Timbuktu, and there is no restriction on the number of intermediate flights. Your query must determine whether a sequence of flights from Madison to Timbuktu exists for *any* input Flights relation instance.

38 Chapter 4

Answer 4.5 In the answers below RA refers to Relational Algebra, TRC refers to Tuple Relational Calculus and DRC refers to Domain Relational Calculus.

1. RA

$$\pi_{eid}(\sigma_{aname='Boeing'}(Aircraft \bowtie Certified))$$

TRC

$$\{C.eid \mid C \in Certified \wedge \\ \exists A \in Aircraft (A.aid = C.aid \wedge A.aname = 'Boeing')\}$$

DRC

$$\{Ceid \mid Ceid, Caid \in Certified \wedge \\ \exists Aid, AN, AR (Aid, AN, AR \in Aircraft \wedge \\ Aid = Caid \wedge AN = 'Boeing')\}$$

SQL

```
SELECT C.eid
FROM Aircraft A, Certified C
WHERE A.aid = C.aid AND A.aname = 'Boeing'
```

2. RA

$$\pi_{ename}(\sigma_{aname='Boeing'}(Aircraft \bowtie Certified \bowtie Employees))$$

TRC

$$\{E.ename \mid E \in Employees \wedge \exists C \in Certified \\ (\exists A \in Aircraft (A.aid = C.aid \wedge A.aname = 'Boeing' \wedge E.eid = C.eid))\}$$

DRC

$$\{EN \mid Eid, EN, ES \in Employees \wedge \\ \exists Ceid, Caid (Ceid, Caid \in Certified \wedge \\ \exists Aid, AN, AR (Aid, AN, AR \in Aircraft \wedge \\ Aid = Caid \wedge AN = 'Boeing' \wedge Eid = Ceid))\}$$

SQL

```
SELECT E.ename
FROM Aircraft A, Certified C, Employees E
WHERE A.aid = C.aid AND A.aname = 'Boeing' AND E.eid = C.eid
```

Relational Algebra and Calculus 39

3. RA

$$\rho(BonnToMadrid, \sigma_{from='Bonn' \wedge to='Madrid'}(Flights)) \\ \pi_{aid}(\sigma_{cruisingrange > distance}(Aircraft \times BonnToMadrid))$$

TRC

$$\{A.aid \mid A \in Aircraft \wedge \exists F \in Flights \\ (F.from = 'Bonn' \wedge F.to = 'Madrid' \wedge A.cruisingrange > F.distance)\}$$

DRC

$$\{Aid \mid Aid, AN, AR \in Aircraft \wedge \\ (\exists FN, FF, FT, FDi, FDe, FA (FN, FF, FT, FDi, FDe, FA \in Flights \wedge \\ FF = 'Bonn' \wedge FT = 'Madrid' \wedge FDi < AR))\}$$

SQL

```
SELECT A.aid
FROM Aircraft A, Flights F
WHERE F.from = 'Bonn' AND F.to = 'Madrid' AND
      A.cruisingrange > F.distance
```

4. RA

$$\pi_{f.lno}(\sigma_{distance < cruisingrange \wedge salary > 100,000}(Flights \bowtie Aircraft \bowtie Certified \bowtie Employees)))$$

TRC $\{F.fno \mid F \in Flights \wedge \exists A \in Aircraft \exists C \in Certified \exists E \in Employees (A.cruisingrange > F.distance \wedge E.salary > 100,000 \wedge A.aid = C.aid \wedge E.eid = C.eid)\}$

DRC

$\{F.N \mid F.N, F.F, F.T, F.Di, F.De, F.A \in Flights \wedge \exists Ceid, Caid (Ceid, Caid \in Certified \wedge \exists Aid, AN, AR (Aid, AN, AR \in Aircraft \wedge \exists Eid, EN, ES (Eid, EN, ES \in Employees (AR > F.Di \wedge ES > 100,000 \wedge Aid = Caid \wedge Eid = Ceid))\}$

SQL

```
SELECT E.ename
FROM Aircraft A, Certified C, Employees E, Flights F
WHERE A.aid = C.aid AND E.eid = C.eid AND
      distance < cruisingrange AND salary > 100,000
```

40 Chapter 4

5. RA $\rho(R1, \pi_{eid}(\sigma_{cruisingrange > 3000}(Aircraft \bowtie Certified))) \pi_{ename}(Employees \bowtie (R1 - \pi_{eid}(\sigma_{aname='Boeing'}(Aircraft \bowtie Certified))))$

TRC

$\{E.ename \mid E \in Employees \wedge \exists C \in Certified (\exists A \in Aircraft (A.aid = C.aid \wedge E.eid = C.eid \wedge A.cruisingrange > 3000)) \wedge \neg (\exists C2 \in Certified (\exists A2 \in Aircraft (A2.aname = 'Boeing' \wedge C2.aid = A2.aid \wedge C2.eid = E.eid)))\}$

DRC

$\{EN \mid Eid, EN, ES \in Employees \wedge \exists Ceid, Caid (Ceid, Caid \in Certified \wedge \exists Aid, AN, AR (Aid, AN, AR \in Aircraft \wedge Aid = Caid \wedge Eid = Ceid \wedge AR > 3000)) \wedge \neg (\exists Aid2, AN2, AR2 (Aid2, AN2, AR2 \in Aircraft \wedge \exists Ceid2, Caid2 (Ceid2, Caid2 \in Certified \wedge Aid2 = Caid2 \wedge Eid = Ceid2 \wedge AN2 = 'Boeing'))\}$

SQL

```
SELECT E.ename
FROM Certified C, Employees E, Aircraft A
WHERE A.aid = C.aid AND E.eid = C.eid AND A.cruisingrange >
```

```

3000 AND E.eid NOT IN ( SELECT C2.eid
FROM Certified C2, Aircraft A2
WHERE C2.aid = A2.aid AND A2.aname = 'Boeing' )

```

6. RA

The approach to take is first find all the employees who do not have the highest salary. Subtract these from the original list of employees and what is left is the highest paid employees.

```

 $\rho(E1, Employees)$ 
 $\rho(E2, Employees)$ 
 $\rho(E3, \pi_{E2.eid}(E1 \text{ } E1.salary > E2.salary \text{ } E2))$ 
 $(\pi_{eid}E1) - E3$ 

```

TRC

```

{E1.eid | E1 ∈ Employees ∧ ¬(∃ E2 ∈ Employees(E2.salary > E1.salary))}

```

DRC

Relational Algebra and Calculus 41

```

{ Eid1 | Eid1, EN1, ES1 ∈ Employees ∧
  ¬(∃ Eid2, EN2, ES2( Eid2, EN2, ES2 ∈ Employees ∧ ES2 > ES1))}

```

SQL

```

SELECT E.eid
FROM Employees E
WHERE E.salary = ( Select MAX (E2.salary)
                  FROM Employees E2 )

```

7. RA

The approach taken is similar to the solution for the previous exercise. First find all the employees who do not have the highest salary. Remove these from the original list of employees and what is left is the highest paid employees. Remove the highest paid employees from the original list. What is left is the second highest paid employees together with the rest of the employees. Then find the highest paid employees of this new list. This is the list of the second highest paid employees.

```

 $\rho(E1, Employees)$ 
 $\rho(E2, Employees)$ 
 $\rho(E3, \pi_{E2.eid}(E1 \text{ } E1.salary > E2.salary \text{ } E2))$ 
 $\rho(E4, E2 \text{ } E3)$ 
 $\rho(E5, E2 \text{ } E3)$ 
 $\rho(E6, \pi_{E5.eid}(E4 \text{ } E1.salary > E5.salary \text{ } E5))$ 
 $(\pi_{eid}E3) - E6$ 

```

TRC

```

{E1.eid | E1 ∈ Employees ∧ ∃ E2 ∈ Employees(E2.salary >
E1.salary ∧ ¬(∃ E3 ∈ Employees(E3.salary > E2.salary)))}

```


DRC

$\{Eid1 \mid Eid1, EN1, ES1 \in Employees \wedge$
 $\exists Eid2, EN2, ES2 (Eid2, EN2, ES2 \in Employees (ES2 > ES1))$
 $\wedge \neg (\exists Eid3, EN3, ES3 (Eid3, EN3, ES3 \in Employees (ES3 > ES2)))\}$

SQL

```
SELECT E.eid
FROM Employees E
WHERE E.salary = (SELECT MAX (E2.salary)
                  FROM Employees E2
                  WHERE E2.salary = (SELECT MAX (E3.salary)
                                      FROM Employees E3 ))
```

42 Chapter 4

8. This cannot be expressed in relational algebra (or calculus) because there is no operator to count, and this query requires the ability to count up to a number that depends on the data. The query can however be expressed in SQL as follows:

```
SELECT Temp.eid
FROM ( SELECT C.eid AS eid, COUNT (C.aid) AS cnt,
        FROM Certified C
        GROUP BY C.eid) AS Temp
WHERE Temp.cnt = ( SELECT MAX (Temp.cnt)
                  FROM Temp)
```

9. RA

The approach behind this query is to first find the employees who are certified for at least three aircraft (they appear at least three times in the Certified relation). Then find the employees who are certified for at least four aircraft. Subtract the second from the first and what is left is the employees who are certified for exactly three aircraft.

$\rho(R1, \text{Certified})$
 $\rho(R2, \text{Certified})$
 $\rho(R3, \text{Certified})$
 $\rho(R4, \text{Certified})$
 $\rho(R5, \pi_{eid}(\sigma_{(R1.eid=R2.eid=R3.eid) \wedge (R1.aid=R2.aid=R3.aid)}(R1 \times R2 \times R3)))$
 $\rho(R6, \pi_{eid}(\sigma_{(R1.eid=R2.eid=R3.eid=R4.eid) \wedge (R1.aid=R2.aid=R3.aid=R4.aid)}(R1 \times R2 \times R3 \times R4)))$
 $R5 - R6$

TRC

$\{C1.eid \mid C1 \in \text{Certified} \wedge \exists C2 \in \text{Certified} (\exists C3 \in \text{Certified}$
 $(C1.eid = C2.eid \wedge C2.eid = C3.eid \wedge$
 $C1.aid = C2.aid \wedge C2.aid = C3.aid \wedge C3.aid = C1.aid \wedge$
 $\neg (\exists C4 \in \text{Certified}$
 $(C3.eid = C4.eid \wedge C1.aid = C4.aid \wedge$
 $C2.aid = C4.aid \wedge C3.aid = C4.aid)))\}$

DRC

$$\begin{aligned}
&\{ CE1 \mid CE1, CA1 \in \text{Certified} \wedge \\
&\exists CE2, CA2 (CE2, CA2 \in \text{Certified} \wedge \\
&\exists CE3, CA3 (CE3, CA3 \in \text{Certified} \wedge \\
&(CE1 = CE2 \wedge CE2 = CE3 \wedge \\
&CA1 = CA2 \wedge CA2 = CA3 \wedge CA3 = CA1 \wedge \\
&\neg (\exists CE4, CA4 (CE4, CA4 \in \text{Certified} \wedge
\end{aligned}$$

Relational Algebra and Calculus 43

$$\begin{aligned}
&(CE3 = CE4 \wedge CA1 = CA4 \wedge \\
&CA2 = CA4 \wedge CA3 = CA4))))\}
\end{aligned}$$

SQL

```

SELECT C1.eid
  FROM Certified C1, Certified C2, Certified C3
 WHERE (C1.eid = C2.eid AND C2.eid = C3.eid AND
        C1.aid = C2.aid AND C2.aid = C3.aid AND C3.aid = C1.aid)
 EXCEPT
 SELECT C4.eid
  FROM Certified C4, Certified C5, Certified C6, Certified C7,
 WHERE (C4.eid = C5.eid AND C5.eid = C6.eid AND C6.eid = C7.eid
        AND C4.aid = C5.aid AND C4.aid = C6.aid AND C4.aid =
        C7.aid AND
        C5.aid = C6.aid AND C5.aid = C7.aid AND C6.aid = C7.aid )

```

This could also be done in SQL using COUNT.

10. This cannot be expressed in relational algebra (or calculus) because there is no operator to sum values. The query can however be expressed in SQL as follows:

```

SELECT SUM (E.salaries)
FROM Employees E

```

11. This cannot be expressed in relational algebra or relational calculus or SQL. The problem is that there is no restriction on the number of intermediate flights. All of the query methods could find if there was a flight directly from Madison to Timbuktu and if there was a sequence of two flights that started in Madison and ended in Timbuktu. They could even find a sequence of n flights that started in Madison and ended in Timbuktu as long as there is a static (i.e., data-independent) upper bound on the number of intermediate flights. (For large n , this would of course be long and impractical, but at least possible.) In this query, however, the upper bound is not static but dynamic (based upon the set of tuples in the Flights relation).

In summary, if we had a static upper bound (say k), we could write an algebra or SQL query that repeatedly computes (upto k) joins on the Flights relation. If the upper bound is dynamic, then we cannot write such a query because k is not known when writing the query.

Exercise 4.6 What is *relational completeness*? If a query language is relationally complete, can you write any desired query in that language?

Answer 4.6 Answer omitted.

Exercise 4.7 What is an *unsafe* query? Give an example and explain why it is important to disallow such queries.

Answer 4.7 An *unsafe* query is a query in relational calculus that has an infinite number of results. An example of such a query is:

$$\{S \mid \neg(S \in \text{Sailors})\}$$

The query is for all things that are not sailors which of course is everything else. Clearly there is an infinite number of answers, and this query is *unsafe*. It is important to disallow *unsafe* queries because we want to be able to get back to users with a list of all the answers to a query after a finite amount of time.

5

SQL: QUERIES, CONSTRAINTS, TRIGGERS

Online material is available for all exercises in this chapter on the book's webpage

at <http://www.cs.wisc.edu/~dbbook>

This includes scripts to create tables for each exercise for use with Oracle, IBM DB2, Microsoft SQL Server, Microsoft Access and MySQL.

Exercise 5.1 Consider the following relations:

Student(*snum*: integer, *sname*: string, *major*: string, *level*: string, *age*: integer)
Class(*name*: string, *meets at*: string, *room*: string, *fid*: integer)
Enrolled(*snum*: integer, *cname*: string)
Faculty(*fid*: integer, *fname*: string, *deptid*: integer)

The meaning of these relations is straightforward; for example, Enrolled has one record per student-class pair such that the student is enrolled in the class.

Write the following queries in SQL. No duplicates should be printed in any of the answers.

1. Find the names of all Juniors (level = JR) who are enrolled in a class taught by I. Teach.

2. Find the age of the oldest student who is either a History major or enrolled in a course taught by I. Teach.
3. Find the names of all classes that either meet in room R128 or have five or more students enrolled.
4. Find the names of all students who are enrolled in two classes that meet at the same time.

45

46 Chapter 5

5. Find the names of faculty members who teach in every room in which some class is taught.
6. Find the names of faculty members for whom the combined enrollment of the courses that they teach is less than five.
7. For each level, print the level and the average age of students for that level.
8. For all levels except JR, print the level and the average age of students for that level.
9. For each faculty member that has taught classes only in room R128, print the faculty member's name and the total number of classes she or he has taught.
10. Find the names of students enrolled in the maximum number of classes.
11. Find the names of students not enrolled in any class.
12. For each age value that appears in Students, find the level value that appears most often. For example, if there are more FR level students aged 18 than SR, JR, or SO students aged 18, you should print the pair (18, FR).

Answer 5.1 The answers are given below:

```
1. SELECT DISTINCT S.Sname
   FROM Student S, Class C, Enrolled E, Faculty F
  WHERE S.snum = E.snum AND E.cname = C.name AND C.fid = F.fid
     AND F.fname = 'I.Teach' AND S.level = 'JR'
```

```
2. SELECT MAX(S.age)
   FROM Student S
  WHERE (S.major = 'History')
     OR S.snum IN (SELECT E.snum
                   FROM Class C, Enrolled E, Faculty F
                  WHERE E.cname = C.name AND C.fid = F.fid
                    AND F.fname = 'I.Teach' )
```

```
3. SELECT C.name
   FROM Class C
  WHERE C.room = 'R128'
     OR C.name IN (SELECT E.cname
                   FROM Enrolled E
                  GROUP BY E.cname
                 HAVING COUNT (*) >= 5)
SQL: Queries, Constraints, Triggers 47
```

```

4. SELECT DISTINCT S.sname
   FROM Student S
   WHERE S.snum IN (SELECT E1.snum
                     FROM Enrolled E1, Enrolled E2, Class C1, Class C2
                     WHERE E1.snum = E2.snum AND E1.cname <> E2.cname
                     AND E1.cname = C1.name
                     AND E2.cname = C2.name AND C1.meets at = C2.meets at)

```

```

5. SELECT DISTINCT F.fname
   FROM Faculty F
   WHERE NOT EXISTS (( SELECT *
                       FROM Class C )
                     EXCEPT
                     (SELECT C1.room
                      FROM Class C1
                      WHERE C1.fid = F.fid ))

```

```

6. SELECT DISTINCT F.fname
   FROM Faculty F
   WHERE 5 > (SELECT COUNT (E.snum)
              FROM Class C, Enrolled E
              WHERE C.name = E.cname
              AND C.fid = F.fid)

```

```

7. SELECT S.level, AVG(S.age)
   FROM Student S
   GROUP BY S.level

```

```

8. SELECT S.level, AVG(S.age)
   FROM Student S
   WHERE S.level <> 'JR'
   GROUP BY S.level

```

```

9. SELECT F.fname, COUNT(*) AS CourseCount
   FROM Faculty F, Class C
   WHERE F.fid = C.fid
   GROUP BY F.fid, F.fname
   HAVING EVERY ( C.room = 'R128' )

```

```

10. SELECT DISTINCT S.sname
    FROM Student S
    WHERE S.snum IN (SELECT E.snum
                     FROM Enrolled E
                     GROUP BY E.snum

```

48 Chapter 5

```

        HAVING COUNT (*) >= ALL (SELECT COUNT (*)
                                FROM Enrolled E2
                                GROUP BY E2.snum ))

```

```

11. SELECT DISTINCT S.sname
    FROM Student S
    WHERE S.snum NOT IN (SELECT E.snum
                        FROM Enrolled E )

```

```

12. SELECT S.age, S.level
      FROM Student S
      GROUP BY S.age, S.level,
      HAVING S.level IN (SELECT S1.level
                        FROM Student S1
                        WHERE S1.age = S.age
                        GROUP BY S1.level, S1.age
                        HAVING COUNT (*) >= ALL (SELECT COUNT (*)
                                              FROM Student S2
                                              WHERE s1.age = S2.age
                                              GROUP BY S2.level, S2.age))

```

Exercise 5.2 Consider the following schema:

```

Suppliers(sid: integer, sname: string, address: string)
Parts(pid: integer, pname: string, color: string)
Catalog(sid: integer, pid: integer, cost: real)

```

The Catalog relation lists the prices charged for parts by Suppliers. Write the following queries in SQL:

1. Find the *pnames* of parts for which there is some supplier.
2. Find the *snames* of suppliers who supply every part.
3. Find the *snames* of suppliers who supply every red part.
4. Find the *pnames* of parts supplied by Acme Widget Suppliers and no one else.
5. Find the *sids* of suppliers who charge more for some part than the average cost of that part (averaged over all the suppliers who supply that part).
6. For each part, find the *sname* of the supplier who charges the most for that part.
7. Find the *sids* of suppliers who supply only red parts.
8. Find the *sids* of suppliers who supply a red part and a green part.

SQL: Queries, Constraints, Triggers 49

9. Find the *sids* of suppliers who supply a red part or a green part.
10. For every supplier that only supplies green parts, print the name of the supplier and the total number of parts that she supplies.
11. For every supplier that supplies a green part and a red part, print the name and price of the most expensive part that she supplies.

Answer 5.2 Answer omitted.

Exercise 5.3 The following relations keep track of airline flight information:

```

Flights(fno: integer, from: string, to: string, distance: integer,
        departs: time, arrives: time, price: real)
Aircraft(aid: integer, aname: string, cruisingrange: integer)
Certified(eid: integer, aid: integer)
Employees(eid: integer, ename: string, salary: integer)

```

Note that the Employees relation describes pilots and other kinds of employees as well; every pilot is certified for some aircraft, and only pilots are certified to fly. Write each of the following queries in SQL. (*Additional queries using the same schema are listed in the exercises for Chapter 4.*)

1. Find the names of aircraft such that all pilots certified to operate them have salaries more than \$80,000.
2. For each pilot who is certified for more than three aircraft, find the *eid* and the maximum *cruisingrange* of the aircraft for which she or he is certified.
3. Find the names of pilots whose *salary* is less than the price of the cheapest route from Los Angeles to Honolulu.
4. For all aircraft with *cruisingrange* over 1000 miles, find the name of the aircraft and the average salary of all pilots certified for this aircraft.
5. Find the names of pilots certified for some Boeing aircraft.
6. Find the *aids* of all aircraft that can be used on routes from Los Angeles to Chicago.
7. Identify the routes that can be piloted by every pilot who makes more than \$100,000.
8. Print the *enames* of pilots who can operate planes with *cruisingrange* greater than 3000 miles but are not certified on any Boeing aircraft.

50 Chapter 5

9. A customer wants to travel from Madison to New York with no more than two changes of flight. List the choice of departure times from Madison if the customer wants to arrive in New York by 6 p.m.
10. Compute the difference between the average salary of a pilot and the average salary of all employees (including pilots).
11. Print the name and salary of every nonpilot whose salary is more than the average salary for pilots.
12. Print the names of employees who are certified only on aircrafts with cruising range longer than 1000 miles.
13. Print the names of employees who are certified only on aircrafts with cruising range longer than 1000 miles, but on at least two such aircrafts.
14. Print the names of employees who are certified only on aircrafts with cruising range longer than 1000 miles and who are certified on some Boeing aircraft.

Answer 5.3 The answers are given below:

1.

```
SELECT DISTINCT A.aname
FROM Aircraft A
WHERE A.Aid IN (SELECT C.aid
                FROM Certified C, Employees E
                WHERE C.eid = E.eid AND
                NOT EXISTS ( SELECT *
                            FROM Employees E1
                            WHERE E1.eid = E.eid AND E1.salary < 80000 ))
```

2. SELECT C.aid, MAX (A.cruisingrange)
 FROM Certified C, Aircraft A
 WHERE C.aid = A.aid
 GROUP BY C.aid
 HAVING COUNT (*) > 3

3. SELECT DISTINCT E.ename
 FROM Employees E
 WHERE E.salary < (SELECT MIN (F.price)
 FROM Flights F
 WHERE F.from = 'Los Angeles' AND F.to = 'Honolulu')

4. Observe that *aid* is the key for Aircraft, but the question asks for aircraft names;
 we deal with this complication by using an intermediate relation Temp:

SQL: Queries, Constraints, Triggers 51

SELECT Temp.name, Temp.AvgSalary
 FROM (SELECT A.aid, A.aname AS name,
 AVG (E.salary) AS AvgSalary
 FROM Aircraft A, Certified C, Employees E
 WHERE A.aid = C.aid AND
 C.aid = E.aid AND A.cruisingrange > 1000
 GROUP BY A.aid, A.aname) AS Temp

5. SELECT DISTINCT E.ename
 FROM Employees E, Certified C, Aircraft A
 WHERE E.aid = C.aid AND
 C.aid = A.aid AND
 A.aname LIKE 'Boeing%'

6. SELECT A.aid
 FROM Aircraft A
 WHERE A.cruisingrange > (SELECT MIN (F.distance)
 FROM Flights F
 WHERE F.from = 'Los Angeles' AND F.to = 'Chicago')

7. SELECT DISTINCT F.from, F.to
 FROM Flights F
 WHERE NOT EXISTS (SELECT *
 FROM Employees E
 WHERE E.salary > 100000
 AND
 NOT EXISTS (SELECT *
 FROM Aircraft A, Certified C
 WHERE A.cruisingrange > F.distance
 AND E.aid = C.aid
 AND A.aid = C.aid))

8. SELECT DISTINCT E.ename
 FROM Employees E
 WHERE E.aid IN ((SELECT C.aid


```

FROM Certified C
WHERE EXISTS ( SELECT A.aid
                FROM Aircraft A
                WHERE A.aid = C.aid
                AND A.cruisingrange > 3000 )
AND
NOT EXISTS ( SELECT A1.aid

```

52 Chapter 5

```

FROM Aircraft A1
WHERE A1.aid = C.aid
AND A1.aname LIKE 'Boeing%' ))

```

```

9. SELECT F.departs
   FROM Flights F
  WHERE F.flno IN ( ( SELECT F0.flno
                     FROM Flights F0
                     WHERE F0.from = 'Madison' AND F0.to = 'New York'
                     AND F0.arrives < '18:00' )
                  UNION
                  ( SELECT F0.flno
                    FROM Flights F0, Flights F1
                    WHERE F0.from = 'Madison' AND F0.to <> 'New York'
                      AND F0.to = F1.from AND F1.to = 'New York'
                      AND F1.departs > F0.arrives
                      AND F1.arrives < '18:00' )
                  UNION
                  ( SELECT F0.flno
                    FROM Flights F0, Flights F1, Flights F2
                    WHERE F0.from = 'Madison'
                      AND F0.to = F1.from
                      AND F1.to = F2.from
                      AND F2.to = 'New York'
                      AND F0.to <> 'New York'
                      AND F1.to <> 'New York'
                      AND F1.departs > F0.arrives
                      AND F2.departs > F1.arrives
                      AND F2.arrives < '18:00' ))

```

```

10. SELECT Temp1.avg - Temp2.avg
     FROM (SELECT AVG (E.salary) AS avg
           FROM Employees E
           WHERE E.eid IN (SELECT DISTINCT C.eid
                          FROM Certified C )) AS Temp1,
          (SELECT AVG (E1.salary) AS avg
           FROM Employees E1 ) AS Temp2

```

```

11. SELECT E.ename, E.salary
     FROM Employees E
    WHERE E.eid NOT IN ( SELECT DISTINCT C.eid
                        FROM Certified C )

```

SQL: Queries, Constraints, Triggers 53

```

AND E.salary > ( SELECT AVG (E1.salary)
                  FROM Employees E1
                  WHERE E1.eid IN
                    ( SELECT DISTINCT C1.eid
                      FROM Certified C1 ) )

```

12. SELECT E.ename
 FROM Employees E, Certified C, Aircraft A
 WHERE C.aid = A.aid AND E.eid = C.eid
 GROUP BY E.eid, E.ename
 HAVING EVERY (A.cruisingrange > 1000)
13. SELECT E.ename
 FROM Employees E, Certified C, Aircraft A
 WHERE C.aid = A.aid AND E.eid = C.eid
 GROUP BY E.eid, E.ename
 HAVING EVERY (A.cruisingrange > 1000) AND COUNT (*) > 1
14. SELECT E.ename
 FROM Employees E, Certified C, Aircraft A
 WHERE C.aid = A.aid AND E.eid = C.eid
 GROUP BY E.eid, E.ename
 HAVING EVERY (A.cruisingrange > 1000) AND ANY (A.aname = 'Boeing')

Exercise 5.4 Consider the following relational schema. An employee can work in more than one department; the *pct time* field of the Works relation shows the percent age of time that a given employee works in a given department.

```

Emp(eid: integer, ename: string, age: integer, salary: real)
Works(eid: integer, did: integer, pct time: integer)
Dept(did: integer, dname: string, budget: real, managerid: integer) Write

```

the following queries in SQL:

1. Print the names and ages of each employee who works in both the Hardware department and the Software department.
2. For each department with more than 20 full-time-equivalent employees (i.e., where the part-time and full-time employees add up to at least that many full-time employees), print the *did* together with the number of employees that work in that department.
3. Print the name of each employee whose salary exceeds the budget of all of the departments that he or she works in.

54 Chapter 5

```

sid sname rating age
18 jones 3 30.0
41 jonah 6 56.0
22 ahab 7 44.0
63 moby null 15.0

```

Figure 5.1 An Instance of Sailors

4. Find the *managerids* of managers who manage only departments with budgets greater than \$1 million.
5. Find the *enames* of managers who manage the departments with the largest budgets.
6. If a manager manages more than one department, he or she *controls* the sum of all the budgets for those departments. Find the *managerids* of managers who control more than \$5 million.
7. Find the *managerids* of managers who control the largest amounts.
8. Find the *enames* of managers who manage only departments with budgets larger than \$1 million, but at least one department with budget less than \$5 million.

Answer 5.4 Answer omitted.

Exercise 5.5 Consider the instance of the Sailors relation shown in Figure 5.1.

1. Write SQL queries to compute the average rating, using AVG; the sum of the ratings, using SUM; and the number of ratings, using COUNT.
2. If you divide the sum just computed by the count, would the result be the same as the average? How would your answer change if these steps were carried out with respect to the *age* field instead of *rating*?
3. Consider the following query: *Find the names of sailors with a higher rating than all sailors with age < 21*. The following two SQL queries attempt to obtain the answer to this question. Do they both compute the result? If not, explain why. Under what conditions would they compute the same result?

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS ( SELECT *
                   FROM Sailors S2
                   WHERE S2.age < 21
                     AND S.rating <= S2.rating )
```

SQL: Queries, Constraints, Triggers 55

```
SELECT *
FROM Sailors S
WHERE S.rating > ANY ( SELECT S2.rating
                     FROM Sailors S2
                     WHERE S2.age < 21 )
```

4. Consider the instance of Sailors shown in Figure 5.1. Let us define instance S1 of Sailors to consist of the first two tuples, instance S2 to be the last two tuples, and S to be the given instance.
 - (a) Show the left outer join of S with itself, with the join condition being *sid=sid*.
 - (b) Show the right outer join of S with itself, with the join condition being *sid=sid*.
 - (c) Show the full outer join of S with itself, with the join condition being *sid=sid*.
 - (d) Show the left outer join of S1 with S2, with the join condition being *sid=sid*.

- (e) Show the right outer join of S1 with S2, with the join condition being *sid=sid*. (f) Show the full outer join of S1 with S2, with the join condition being *sid=sid*.

Answer 5.5 The answers are shown below:

1. SELECT AVG (S.rating) AS AVERAGE
FROM Sailors S

SELECT SUM (S.rating)
FROM Sailors S

SELECT COUNT (S.rating)
FROM Sailors S

2. The result using SUM and COUNT would be smaller than the result using AVERAGE if there are tuples with rating = NULL. This is because all the aggregate operators, except for COUNT, ignore NULL values. So the first approach would compute the average over all tuples while the second approach would compute the average over all tuples with non-NULL rating values. However, if the aggregation is done on the age field, the answers using both approaches would be the same since the age field does not take NULL values.
3. Only the first query is correct. The second query returns the names of sailors with a higher rating than *at least one* sailor with age < 21. Note that the answer to the second query does not necessarily contain the answer to the first query. In particular, if all the sailors are at least 21 years old, the second query will return an empty set while the first query will return all the sailors. This is because the NOT EXISTS predicate in the first query will evaluate to *true* if its subquery evaluates

56 Chapter 5

	<i>sid sname rating age</i>	<i>sid sname rating age</i>
4. (a) (b)	18 jones 3 30.0	18 jones 3 30.0
	rating age 18 jones 3 30.0	18
	jones 3 30.0 41 jonah 6 56.0	41
	jonah 6 56.0 22 ahab 7 44.0	22
	ahab 7 44.0 63 moby null 15.0	63
	moby null 15.0	
(c)	<i>sid sname rating age</i>	<i>sid sname</i>
	rating age 18 jones 3 30.0	18
41 jonah 6 56.0	41 jonah 6 56.0	41
22 ahab 7 44.0	22 ahab 7 44.0	22
63 moby null 15.0	63 moby null 15.0	63
	moby null 15.0	
	<i>sid sname rating age</i>	<i>sid sname</i>

to an empty set, while the ANY predicate in the second query will evaluate to *false* if its subquery evaluates to an empty set. The two queries give the same results if and only if one of the following two conditions hold:

The *Sailors* relation is empty, or

There is at least one sailor with age > 21 in the *Sailors* relation, and for

every sailor *s*, either *s* has a higher rating than all sailors under 21 or *s* has a rating no higher than all sailors under 21.

Exercise 5.6 Answer the following questions:

1. Explain the term *impedance mismatch* in the context of embedding SQL commands in a host language such as C.

sid sname rating age sid sname rating age

(d)

18 jones 3 30.0 null null null null 41 jonah 6 56.0 null null
null null

sid sname rating age sid sname rating age (e)

null null null null 22 ahab 7 44.0 null null null null 63 moby
null 15.0

SQL: Queries, Constraints, Triggers 57

sid sname rating age sid sname rating age

18 jones 3 30.0 null null null null

(f)

41 jonah 6 56.0 null null null null null null null null 22 ahab 7
44.0 null null null null 63 moby null 15.0

2. How can the value of a host language variable be passed to an embedded SQL command?
3. Explain the WHENEVER command's use in error and exception handling.
4. Explain the need for cursors.
5. Give an example of a situation that calls for the use of embedded SQL; that is, in interactive use of SQL commands is not enough, and some host language capabilities are needed.
6. Write a C program with embedded SQL commands to address your example in the previous answer.
7. Write a C program with embedded SQL commands to find the standard deviation of sailors' ages.
8. Extend the previous program to find all sailors whose age is within one standard deviation of the average age of all sailors.
9. Explain how you would write a C program to compute the transitive closure of a graph, represented as an SQL relation Edges(*from*, *to*), using embedded SQL commands. (You need not write the program, just explain the main points to be dealt with.)
10. Explain the following terms with respect to cursors: *updatability*, *sensitivity*, and *scrollability*.
11. Define a cursor on the Sailors relation that is updatable, scrollable, and returns answers sorted by *age*. Which fields of Sailors can such a cursor *not* update? Why?

12. Give an example of a situation that calls for dynamic SQL; that is, even embedded SQL is not sufficient.

Answer 5.6 Answer omitted.

Exercise 5.7 Consider the following relational schema and briefly answer the questions that follow:

58 Chapter 5

Emp(*eid*: integer, *ename*: string, *age*: integer, *salary*: real)
Works(*eid*: integer, *did*: integer, *pct time*: integer)
Dept(*did*: integer, *budget*: real, *managerid*: integer)

1. Define a table constraint on Emp that will ensure that every employee makes at least \$10,000.
2. Define a table constraint on Dept that will ensure that all managers have *age* > 30.
3. Define an assertion on Dept that will ensure that all managers have *age* > 30. Compare this assertion with the equivalent table constraint. Explain which is better.
4. Write SQL statements to delete all information about employees whose salaries exceed that of the manager of one or more departments that they work in. Be sure to ensure that all the relevant integrity constraints are satisfied after your updates.

Answer 5.7 The answers are given below:

1. Define a table constraint on Emp that will ensure that every employee makes at least \$10,000

```
CREATE TABLE Emp ( eid INTEGER,
                    ename CHAR(10),
                    age INTEGER ,
                    salary REAL,
                    PRIMARY KEY (eid),
                    CHECK ( salary >= 10000 ))
```

2. Define a table constraint on Dept that will ensure that all managers have *age* > 30

```
CREATE TABLE Dept ( did INTEGER,
                    buget REAL,
                    managerid INTEGER ,
                    PRIMARY KEY (did),
                    FOREIGN KEY (managerid) REFERENCES Emp,
                    CHECK ( ( SELECT E.age FROM Emp E, Dept D
                              WHERE E.eid = D.managerid ) > 30 )
```

3. Define an assertion on Dept that will ensure that all managers have *age* > 30

```
CREATE TABLE Dept ( did INTEGER,
                    budget REAL,
                    managerid INTEGER ,
                    PRIMARY KEY (did) )
```

SQL: Queries, Constraints, Triggers 59

```

CREATE ASSERTION managerAge
CHECK ( (SELECT E.age
        FROM Emp E, Dept D
        WHERE E.eid = D.managerid ) > 30 )

```

Since the constraint involves two relations, it is better to define it as an assertion, independent of any one relation, rather than as a check condition on the Dept relation. The limitation of the latter approach is that the condition is checked only when the Dept relation is being updated. However, since age is an attribute of the Emp relation, it is possible to update the age of a manager which violates the constraint. So the former approach is better since it checks for potential violation of the assertion whenever one of the relations is updated.

4. To write such statements, it is necessary to consider the constraints defined over the tables. We will assume the following:

```

CREATE TABLE Emp ( eid INTEGER,
                    ename CHAR(10),
                    age INTEGER,
                    salary REAL,
                    PRIMARY KEY (eid) )

CREATE TABLE Works ( eid INTEGER,
                      did INTEGER,
                      pcttime INTEGER,
                      PRIMARY KEY (eid, did),
                      FOREIGN KEY (did) REFERENCES Dept,
                      FOREIGN KEY (eid) REFERENCES Emp,
                      ON DELETE CASCADE)

CREATE TABLE Dept ( did INTEGER,
                     budget REAL,
                     managerid INTEGER ,
                     PRIMARY KEY (did),
                     FOREIGN KEY (managerid) REFERENCES Emp,
                     ON DELETE SET NULL)

```

Now, we can define statements to delete employees who make more than one of their managers:

```

DELETE
FROM Emp E
WHERE E.eid IN ( SELECT W.eid
                 FROM Work W, Emp E2, Dept D
                 WHERE W.did = D.did

```

60 Chapter 5

```

                 AND D.managerid = E2.eid
                 AND E.salary > E2.salary )

```

Exercise 5.8 Consider the following relations:

Student(*snum*: integer, *sname*: string, *major*: string,

level: string, *age*: integer)
 Class(*name*: string, *meets at*: time, *room*: string, *fid*: integer)
 Enrolled(*snum*: integer, *cname*: string)
 Faculty(*fid*: integer, *fname*: string, *deptid*: integer)

The meaning of these relations is straightforward; for example, Enrolled has one record per student-class pair such that the student is enrolled in the class.

1. Write the SQL statements required to create these relations, including appropriate versions of all primary and foreign key integrity constraints.
2. Express each of the following integrity constraints in SQL unless it is implied by the primary and foreign key constraint; if so, explain how it is implied. If the constraint cannot be expressed in SQL, say so. For each constraint, state what operations (inserts, deletes, and updates on specific relations) must be monitored to enforce the constraint.
 - (a) Every class has a minimum enrollment of 5 students and a maximum enrollment of 30 students.
 - (b) At least one class meets in each room.
 - (c) Every faculty member must teach at least two courses.
 - (d) Only faculty in the department with *deptid*=33 teach more than three courses.
 - (e) Every student must be enrolled in the course called Math101.
 - (f) The room in which the earliest scheduled class (i.e., the class with the smallest *meets at* value) meets should not be the same as the room in which the latest scheduled class meets.
 - (g) Two classes cannot meet in the same room at the same time.
 - (h) The department with the most faculty members must have fewer than twice the number of faculty members in the department with the fewest faculty members.
 - (i) No department can have more than 10 faculty members.
 - (j) A student cannot add more than two courses at a time (i.e., in a single update).
 - (k) The number of CS majors must be more than the number of Math majors.

SQL: Queries, Constraints, Triggers 61

- (l) The number of distinct courses in which CS majors are enrolled is greater than the number of distinct courses in which Math majors are enrolled.
- (m) The total enrollment in courses taught by faculty in the department with *deptid*=33 is greater than the number of Math majors.
- (n) There must be at least one CS major if there are any students whatsoever.
- (o) Faculty members from different departments cannot teach in the same room.

Answer 5.8 Answer omitted.

Exercise 5.9 Discuss the strengths and weaknesses of the trigger mechanism. Contrast triggers with other integrity constraints supported by SQL.

Answer 5.9 A trigger is a procedure that is automatically invoked in response to a specified change to the database. The advantages of the trigger mechanism include

the ability to perform an action based on the result of a query condition. The set of actions that can be taken is a superset of the actions that integrity constraints can take (i.e. report an error). Actions can include invoking new update, delete, or insert queries, perform data definition statements to create new tables or views, or alter security policies. Triggers can also be executed before or after a change is made to the database (that is, use old or new data).

There are also disadvantages to triggers. These include the added complexity when trying to match database modifications to trigger events. Also, integrity constraints are incorporated into database performance optimization; it is more difficult for a database to perform automatic optimization with triggers. If database consistency is the primary goal, then integrity constraints offer the same power as triggers. Integrity constraints are often easier to understand than triggers.

Exercise 5.10 Consider the following relational schema. An employee can work in more than one department; the *pct time* field of the Works relation shows the percent age of time that a given employee works in a given department.

```
Emp(eid: integer, ename: string, age: integer, salary: real)
Works(eid: integer, did: integer, pct time: integer)
Dept(did: integer, budget: real, managerid: integer)
```

Write SQL-92 integrity constraints (domain, key, foreign key, or CHECK constraints; or assertions) or SQL:1999 triggers to ensure each of the following requirements, considered independently.

1. Employees must make a minimum salary of \$1000.

62 Chapter 5

2. Every manager must be also be an employee.

3. The total percentage of all appointments for an employee must be under 100%.

4. A manager must always have a higher salary than any employee that he or she manages.

5. Whenever an employee is given a raise, the manager's salary must be increased to be at least as much.

6. Whenever an employee is given a raise, the manager's salary must be increased to be at least as much. Further, whenever an employee is given a raise, the department's budget must be increased to be greater than the sum of salaries of all employees in the department.

Answer 5.10 Answer omitted.

6

DATABASE APPLICATION DEVELOPMENT

Exercise 6.1 Briefly answer the following questions.

1. Explain the following terms: Cursor, Embedded SQL, JDBC, SQLJ, stored procedure.
2. What are the differences between JDBC and SQLJ? Why do they both exist?
3. Explain the term *stored procedure*, and give examples why stored procedures are useful.

Answer 6.1 The answers are given below:

1. A cursor enables individual row access of a relation by positioning itself at a row and reading its contents. Embedded SQL refers to the usage of SQL commands within a host program. JDBC stands for Java DataBase Connectivity and is an interface that allows a Java program to easily connect to any database system. SQLJ is a tool that allows SQL to be embedded directly into a Java program. A stored procedure is program that runs on the database server and can be called with a single SQL statement.
2. SQLJ provides embedded SQL statements. These SQL statements are static in nature and thus are preprocessed and precompiled. For instance, syntax checking and schema checking are done at compile time. JDBC allows dynamic queries that are checked at runtime. SQLJ is easier to use than JDBC and is often a better option for static queries. For dynamic queries, JDBC must still be used.
3. Stored procedures are programs that run on the database server and can be called with a single SQL statement. They are useful in situations where the processing should be done on the server side rather than the client side. Also, since the procedures are centralized to the server, code writing and maintenance is simplified, because the client programs do not have to duplicate the application logic. Stored procedures can also be used to reduce network communication; the results of a stored procedure can be analyzed and kept on the database server.

Exercise 6.2 Explain how the following steps are performed in JDBC:

1. Connect to a data source.
2. Start, commit, and abort transactions.
3. Call a stored procedure.

How are these steps performed in SQLJ?

Answer 6.2 The answers are given below:

1. Connecting to a data source in JDBC involves the creation of a *Connection* object. Parameters for the connection are specified using a *JDBC URL* that contains things like the network address of the database server and the username and password for connecting.
SQLJ makes calls to the same JDBC driver for connecting to a data source and uses the same type of JDBC URL.
2. Each connection can specify how to handle transactions. If the *autocommit* flag is set, each SQL statement is treated as a separate transaction. If the flag is turned off, there is a *commit()* function call that will actually commit the transaction. The autocommit flag can also be set in SQLJ. If the flag is not set, transactions are committed by passing a *COMMIT* SQL statement to the DBMS.
3. Stored procedures are called from JDBC using the *CallableStatement* class with the SQL command *{CALL StoredProcedureName}*.
SQLJ also uses *CALL StoredProcedureName* to execute stored procedures at the DBMS.

Exercise 6.3 Compare exception handling and handling of warnings in embedded SQL, dynamic SQL, JDBC, and SQLJ.

Answer 6.3 The answers are given below:

Embedded SQL: The *SQLSTATE* variable is used to check for errors after each Embedded SQL statement is executed. If an error has occurred, program control is transferred to a separate statement. This is done during the precompilation step for static queries.

Dynamic SQL: For dynamic SQL, the SQL statement can change at runtime and thus the error handling must also occur at runtime.

JDBC: In JDBC, programmers can use the *try ... catch* syntax to handle exceptions of type *SQLException*. The *SQLWarning* class is used for problems not as severe as errors. They are not caught in the *try ... catch* statement and must be checked independently with a *getWarnings()* function call.

Database Application Development 65 SQLJ: SQLJ uses the same

mechanisms as JDBC to catch error and warnings.

Exercise 6.4 Answer the following questions.

1. Why do we need a precompiler to translate embedded SQL and SQLJ? Why do we not need a precompiler for JDBC?
2. SQLJ and embedded SQL use variables in the host language to pass parameters to SQL queries, whereas JDBC uses placeholders marked with a '?'. Explain the difference, and why the different mechanisms are needed.

Answer 6.4 Answer omitted.

Exercise 6.5 A dynamic web site generates HTML pages from information stored in a database. Whenever a page is requested, is it dynamically assembled from static

data and data in a database, resulting in a database access. Connecting to the database is usually a time-consuming process, since resources need to be allocated, and the user needs to be authenticated. Therefore, connection pooling—setting up a pool of persistent database connections and then reusing them for different requests can significantly improve the performance of database-backed websites. Since servlets can keep information beyond single requests, we can create a connection pool, and allocate resources from it to new requests.

Write a connection pool class that provides the following methods:

Create the pool with a specified number of open connections to the database system.

Obtain an open connection from the pool.

Release a connection to the pool.

Destroy the pool and close all connections.

Answer 6.5 The answer for this exercise is available online for instructors. To find out how to get access to instructor's material, visit the book homepage at <http://www.cs.wisc.edu/~dbbook>.

7

INTERNET APPLICATIONS

Exercise 7.1 Briefly answer the following questions:

1. Explain the following terms and describe what they are used for: HTML, URL, XML, Java, JSP, XSL, XSLT, servlet, cookie, HTTP, CSS, DTD.
2. What is CGI? Why was CGI introduced? What are the disadvantages of an architecture using CGI scripts?
3. What is the difference between a webserver and an application server? What functionality do typical application servers provide?
4. When is an XML document well-formed? When is an XML document valid?

Answer 7.1 The answers are as follows.

1. *HTTP* (HyperText Transfer Protocol) is the communication protocol used to con

nect clients with servers over the Internet. *URL* (Universal Resource Locator) is a string that uniquely identifies an internet address. *HTML* (HyperText Markup Language) is a simple language used to enhance regular text by including special tags. *CSS* (Cascading Style Sheets) are used to define how to display HTML documents. *XML* (Extensible Markup Language) allows users to define their own markup tags in a document. *XSL* (Extensible Style Language) can be used to describe how an XML document should be displayed. *XSLT* (XML Transformation Language) is a language that can transform input XML into differently structured XML. A *DTD* (Document Type Declaration) is a grammar that describes how to use new tags in an XML document. *Java* is cross-platform interpreted programming language. *Servlets* are pieces of Java code that run on the middle tier or server layers and be used for any functionality that Java provides. *JSP* (JavaServer Pages) are HTML pages with embedded servlet code.. *Cookies* are a simple way to store persistent data at the client level.

Figure 7.1 Solution to Exercise 7.2 (d)

2. *CGI* (Common Gateway Interface) specifies how the web server communicates other programs on the server. CGI programs are used to pass HTML form data to other programs that process that data. Each page request will create a new process on the server, which is a performance issue when requests are scaled up.
3. A web server handles the interaction with the client's web browser. Application servers are used to maintain a pool of processes for handling requests. Typically, they are the middleware tier between the web server and the data sources such as database systems. Application servers eliminate the problems with process creation overload and can also provide extra functionality like abstracting away heterogeneous data sources and maintaining session state information.
4. An XML document is *valid* if it has an associated DTD and the document follows the rules of the DTD. An XML document is *well-formed* if it follows three guide lines: (1) it starts with an XML declaration, (2) it contains a root element that contains all other elements and (3) all elements are properly nested.

Exercise 7.2 Briefly answer the following questions about the HTTP protocol:

1. What is a communication protocol?
2. What is the structure of an HTTP request message? What is the structure of an HTTP response message? Why do HTTP messages carry a version field?
3. What is a stateless protocol? Why was HTTP designed to be stateless?
4. Show the HTTP request message generated when you request the home page of this book (<http://www.cs.wisc.edu/~dbbook>). Show the HTTP response message that the server generates for that page.

Answer 7.2 Answer omitted.

Exercise 7.3 In this exercise, you are asked to write the functionality of a generic shopping basket; you will use this in several subsequent project exercises. Write a set of JSP pages that displays a shopping basket of items and allows users to add, remove, and change the quantity of items. To do this, use a cookie storage scheme that stores the following information:

The UserId of the user who owns the shopping basket.

The number of products stored in the shopping basket.

68 Chapter 7 A product id and a quantity for each product.

When manipulating cookies, remember to set the Expires property such that the cookie can persist for a session or indefinitely. Experiment with cookies using JSP and make sure you know how to retrieve, set values, and delete the cookie.

You need to create five JSP pages to make your prototype complete:

Index Page (index.jsp): This is the main entry point. It has a link that directs the user to the Products page so they can start shopping.

Products Page (products.jsp): Shows a listing of all products in the database with their descriptions and prices. This is the main page where the user fills out the shopping basket. Each listed product should have a button next to it, which adds it to the shopping basket. (If the item is already in the shopping basket, it increments the quantity by one.) There should also be a counter to show the total number of items currently in the shopping basket. Note that if a user has a quantity of five of a single item in the shopping basket, the counter should indicate a total quantity of five. The page also contains a button that directs the user to the Cart page.

Cart Page (cart.jsp): Shows a listing of all items in the shopping basket cookie. The listing for each item should include the product name, price, a text box for the quantity (the user can change the quantity of items here), and a button to remove the item from the shopping basket. This page has three other buttons: one button to continue shopping (which returns the user to the Products page), a second button to update the cookie with the altered quantities from the text boxes, and a third button to place or confirm the order, which directs the user to the Confirm page.

Confirm Page (confirm.jsp): Lists the final order. There are two buttons on this page. One button cancels the order and the other submits the completed order. The cancel button just deletes the cookie and returns the user to the Index page. The submit button updates the database with the new order, deletes the cookie, and returns the user to the Index page.

Exercise 7.4 In the previous exercise, replace the page products.jsp with the following *search page* search.jsp. This page allows users to search products by name or description. There should be both a text box for the search text and radio buttons to allow the user to choose between search-by-name and search-by-description (as well as a submit button to retrieve the results). The page that handles search results should be modeled after products.jsp (as described in the previous exercise) and be called products.jsp. It should retrieve all records where the search text is a substring of the name or description (as chosen by the user). To integrate this with the previous exercise, simply replace all the links to products.jsp with search.jsp.

Exercise 7.5 Write a simple authentication mechanism (without using encrypted transfer of passwords, for simplicity). We say a user is authenticated if she has provided a valid username-password combination to the system; otherwise, we say the user is not authenticated. Assume for simplicity that you have a database schema that stores only a customer id and a password:

Passwords(cid: integer, username: string, password: string)

1. How and where are you going to track when a user is 'logged on' to the system?
2. Design a page that allows a registered user to log on to the system.
3. Design a page header that checks whether the user visiting this page is logged in.

Exercise 7.6 (Due to Jeff Derstadt) TechnoBooks.com is in the process of reorganizing its website. A major issue is how to efficiently handle a large number of search results. In a human interaction study, it found that modem users typically like to view 20 search results at a time, and it would like to program this logic into the system. Queries that return batches of sorted results are called *top N queries*. (See Section 23 for a discussion of database support for top N queries.) For example, results 1-20 are returned, then results 21-40, then 41-60, and so on. Different techniques are used for performing top N queries and TechnoBooks.com would like you to implement two of them.

Infrastructure: Create a database with a table called Books and populate it with some books, using the format that follows. This gives you 111 books in your database with a title of AAA, BBB, CCC, DDD, or EEE, but the keys are not sequential for books with the same title.

Books(bookid: INTEGER, title: CHAR(80), author: CHAR(80), price: REAL)

```

For i = 1 to 111 {
    Insert the tuple (i, "AAA", "AAA Author", 5.99)
    i=i+1
    Insert the tuple (i, "BBB", "BBB Author", 5.99)
    i=i+1
    Insert the tuple (i, "CCC", "CCC Author", 5.99)
    i=i+1
    Insert the tuple (i, "DDD", "DDD Author", 5.99)
    i=i+1
    Insert the tuple (i, "EEE", "EEE Author", 5.99)
}

```

70 Chapter 7

Placeholder Technique: The simplest approach to top N queries is to store a placeholder for the first and last result tuples, and then perform the same query. When the new query results are returned, you can iterate to the placeholders and return the previous or next 20 results.

Tuples Shown Lower Placeholder Previous Set Upper Placeholder Next Set 1-20
 1 None 20 21-40 21-40 21 1-20 40 41-60 41-60 41 21-40 60 61-80

Write a webpage in JSP that displays the contents of the Books table, sorted by the Title and BookId, and showing the results 20 at a time. There should be a link (where appropriate) to get the previous 20 results or the next 20 results. To do this, you can encode the placeholders in the Previous or Next Links as follows. Assume that you are displaying records 21–40. Then the previous link is `display.jsp?lower=21` and the next link is `display.jsp?upper=40`.

You should not display a previous link when there are no previous results; nor should you show a Next link if there are no more results. When your page is called again to get another batch of results, you can perform the same query to get all the records, iterate through the result set until you are at the proper starting point, then display 20 more results.

What are the advantages and disadvantages of this technique?

Query Constraints Technique: A second technique for performing top N queries is to push boundary constraints into the query (in the WHERE clause) so that the query returns only results that have not yet been displayed. Although this changes the query, fewer results are returned and it saves the cost of iterating up to the boundary. For example, consider the following table, sorted by (title, primary key).

Internet Applications 71

Batch	Result Number	Title	Primary Key
	1	AAA	105
	2	BBB	13
	3	CCC	48
	4	DDD	52
	5	DDD	101
2	6	DDD	121
	7	EEE	19
	8	EEE	68
	9	FFF	2
2	10	FFF	33
	11	FFF	58
	12	FFF	59
3	13	GGG	93
	14	HHH	132
	15	HHH	135

In batch 1, rows 1 through 5 are displayed, in batch 2 rows 6 through 10 are displayed, and so on. Using the placeholder technique, all 15 results would be returned for each batch. Using the constraint technique, batch 1 displays results 1-5 but returns results 1-15, batch 2 will display results 6-10 but returns only results 6-15, and batch 3 will display results 11-15 but return only results 11-15.

The constraint can be pushed into the query because of the sorting of this table. Consider the following query for batch 2 (displaying results 6-10):

```
EXEC SQL SELECT B.Title
FROM Books B
WHERE (B.Title = 'DDD' AND B.BookId > 101) OR (B.Title > 'DDD')
ORDER BY B.Title, B.BookId
```


This query first selects all books with the title 'DDD,' but with a primary key that is greater than that of record 5 (record 5 has a primary key of 101). This returns record 6. Also, any book that has a title after 'DDD' alphabetically is returned. You can then display the first five results.

The following information needs to be retained to have Previous and Next buttons that return more results:

Previous: The title of the *first* record in the previous set, and the primary key of the *first* record in the previous set.

Next: The title of the *first* record in the next set; the primary key of the *first* record in the next set.

72 Chapter 7

These four pieces of information can be encoded into the Previous and Next buttons as in the previous part. Using your database table from the first part, write a JavaServer Page that displays the book information 20 records at a time. The page should include *Previous* and *Next* buttons to show the previous or next record set if there is one. Use the constraint query to get the Previous and Next record sets.

8

OVERVIEW OF STORAGE AND INDEXING

Exercise 8.1 Answer the following questions about data on external storage in a DBMS:

1. Why does a DBMS store data on external storage?
2. Why are I/O costs important in a DBMS?
3. What is a record id? Given a record's id, how many I/Os are needed to fetch it into main memory?
4. What is the role of the buffer manager in a DBMS? What is the role of the disk space manager? How do these layers interact with the file and access methods layer?

Answer 8.1 The answer to each question is given below.

1. A DBMS stores data on external storage because the quantity of data is vast,

and must persist across program executions.

2. I/O costs are of primary importance to a DBMS because these costs typically dominate the time it takes to run most database operations. Optimizing the amount of I/O's for an operation can result in a substantial increase in speed in the time it takes to run that operation.
3. A record id, or rid for short, is a unique identifier for a particular record in a set of records. An rid has the property that we can identify the disk address of the page containing the record by using the rid. The number of I/O's required to read a record, given a rid, is therefore 1 I/O.
4. In a DBMS, the buffer manager reads data from persistent storage into memory as well as writes data from memory into persistent storage. The disk space manager manages the available physical storage space of data for the DBMS. When the file

and access methods layer needs to process a page, it asks the buffer manager to fetch the page and put it into memory if it is not already in memory. When the files and access methods layer needs additional space to hold new records in a file, it asks the disk space manager to allocate an additional disk page.

Exercise 8.2 Answer the following questions about files and indexes:

1. What operations are supported by the file of records abstraction?
2. What is an index on a file of records? What is a search key for an index? Why do we need indexes?
3. What alternatives are available for the data entries in an index?
4. What is the difference between a primary index and a secondary index? What is a duplicate data entry in an index? Can a primary index contain duplicates?
5. What is the difference between a clustered index and an unclustered index? If an index contains data records as 'data entries,' can it be unclustered?
6. How many clustered indexes can you create on a file? Would you always create at least one clustered index for a file?
7. Consider Alternatives (1), (2) and (3) for 'data entries' in an index, as discussed in Section 8.2. Are all of them suitable for secondary indexes? Explain.

Answer 8.2 Answer omitted.

Exercise 8.3 Consider a relation stored as a randomly ordered file for which the only index is an unclustered index on a field called *sal*. If you want to retrieve all records with *sal* > 20, is using the index always the best alternative? Explain.

Answer 8.3 No. In this case, the index is unclustered, each qualifying data entry could contain an rid that points to a distinct data page, leading to as many data page I/Os as the number of data entries that match the range query. In this situation, using index is actually worse than file scan.

Exercise 8.4 Consider the instance of the Students relation shown in Figure 8.1,

sorted by *age*: For the purposes of this question, assume that these tuples are stored in a sorted file in the order shown; the first tuple is on page 1 the second tuple is also on page 1; and so on. Each page can store up to three data records; so the fourth tuple is on page 2.

Explain what the data entries in each of the following indexes contain. If the order of entries is significant, say so and explain why. If such an index cannot be constructed, say so and explain why.

Overview of Storage and Indexing 75

	<i>sid</i>	<i>name</i>	<i>login</i>	<i>age</i>	<i>gpa</i>
	53831	Madayan	madayan@music	11	1.8
	53832	Guldu	guldu@music	12	2.0
	53666	Jones	jones@cs	18	3.4
	53688	Smith	smith@ee	19	3.2
	53650	Smith	smith@math	19	3.8

Figure 8.1 An Instance of the Students Relation, Sorted by *age*

1. An unclustered index on *age* using Alternative (1).
2. An unclustered index on *age* using Alternative (2).
3. An unclustered index on *age* using Alternative (3).
4. A clustered index on *age* using Alternative (1).
5. A clustered index on *age* using Alternative (2).
6. A clustered index on *age* using Alternative (3).
7. An unclustered index on *gpa* using Alternative (1).
8. An unclustered index on *gpa* using Alternative (2).
9. An unclustered index on *gpa* using Alternative (3).
10. A clustered index on *gpa* using Alternative (1).
11. A clustered index on *gpa* using Alternative (2).
12. A clustered index on *gpa* using Alternative (3).

Answer 8.4 Answer omitted.

Exercise 8.5 Explain the difference between Hash indexes and B+-tree indexes. In particular, discuss how equality and range searches work, using an example.

Answer 8.5 A Hash index is constructed by using a hashing function that quickly maps a search key value to a specific location in an array-like list of elements called buckets. The buckets are often constructed such that there are more bucket locations than there are possible search key values, and the hashing function is chosen so that it is not often that two search key values hash to the same bucket. A B+-tree index is constructed by sorting the data on the search key and maintaining a hierarchical search data structure that directs searches to the correct page of data entries.